

# ECE 570/670

David Irwin  
Lecture 3

# Administrative Details

- Website: <https://courses.umass.edu/eceng570-deirwin/spring23>
  - Username: **irwin-ece**, Password: **myclass**
  - *Calendar will change*
- Think about Assignment partners
  - Assignment 0 will be posted on Thursday
    - Fill out Google Form your NetID and NetID of your partner
      - <https://forms.gle/ChgjQehA5hQku5W6>
      - I will send back to you a username (“groupXX”) and a password to login to the VM you **must** use for the assignments
    - **This must be done by next Sunday!**
  - Ensure you have and can check **@umass.edu** email
- UNIX paper posted for Thursday

# Administrative Details

- Piazza setup
  - Will update tomorrow after grad add/drop
- Reviewing system
  - Will add people tomorrow after grad add/drop
  - Link to reviewing system on website
- Moodle setup
  - Only used for posting grades
- VM image posted (ECE670-Spring23.ova)
  - Website resources page
- My Office hours
  - Mondays 2:30pm – 3:30pm
  - Thursdays 9am – 10am

# Last Time

- Finished intro
- Started talking about processes and threads

# Today's Outline

- Discuss paper
- Continue talking about threads
- Start talking about how to deal with synchronization issues with multiple threads
  - Thread cooperation

# “Hints for Computer System Design”

- Technical fellow at Microsoft
- Adjunct Professor at MIT
- Computer architecture, LANs, printers, operating systems, RPCs, programming languages, security, etc.
- Won Turing award in 1992
- Questions? Comments?
- Discussion.



Butler Lampson

# “Hints for Computer System Design”

- Divides Hints into 3 categories
  - 1. Functionality
  - 2. Speed
  - 3. Fault-tolerance

# Functionality

- Keep it Simple
  - “Do one thing and do it well”
  - “Get it right”
  - “Don’t hide power”
  - “Leave it to the client”
- Continuity
  - “Keep interfaces stable”
  - “Keep a place to stand”
- Making things work
  - “Plan to throw one away”
  - “Keep Secrets”
  - “Divide and Conquer”

# Speed

- Suggestions
  - Think about “splitting resources” rather than sharing
  - “Cache answers” or use hints
  - “When in doubt use brute force”
  - “Compute in the background”
  - “Use batch processing if possible”
  - “Safety first”
  - “Shed load”

# Fault-tolerance

- Suggestions
  - Use end-to-end error recovery
  - Log updates before making them
  - Make actions either atomic or restartable
    - What's another word for restartable?

# Consider a web server

- One processor
- Multiple disks
- Tasks
  - 1. Receives multiple, simultaneous requests
    - *Return from accept() with socket ids of new connections*
  - 2. Reads web pages from disk
    - *Receive data by calling recv() on socket id and parse request to find out request HTML web page*
  - 3. Returns on-disk files to requester
    - *Call read() to read the requested HTML file from the file system, and call send() to send the data to the requester*

# Three Different Designs

- ***Single-threaded*** – one request at a time, lots of waiting
- ***Event-driven*** - one thread, but interleave requests
- ***Multi-threaded*** – many threads, one request per thread, lots of waiting but scheduler switches between threads

# Benefits of Threads

- Thread manager takes care of sharing CPU among threads
  - Thread can issue blocking I/Os, while other threads make progress
  - Private state for each thread
- Applications get a simpler programming model
  - Illusion of a dedicated CPU per thread
- Can isolate requests and schedule them differently
- Any drawbacks?
  - Overhead of context switching (matters in high concurrency servers)
  - Possible state overhead (replicated state between threads)

# Threads are useful

- They cannot provide total independence
  - But they are still a useful abstraction!
  - Threads make concurrent programming easier
- Thread system manages sharing the CPU
  - (unlike in the event-driven case)
- Apps can *encapsulate* task state within a thread
  - (e.g. web request state)

# Where are threads used?

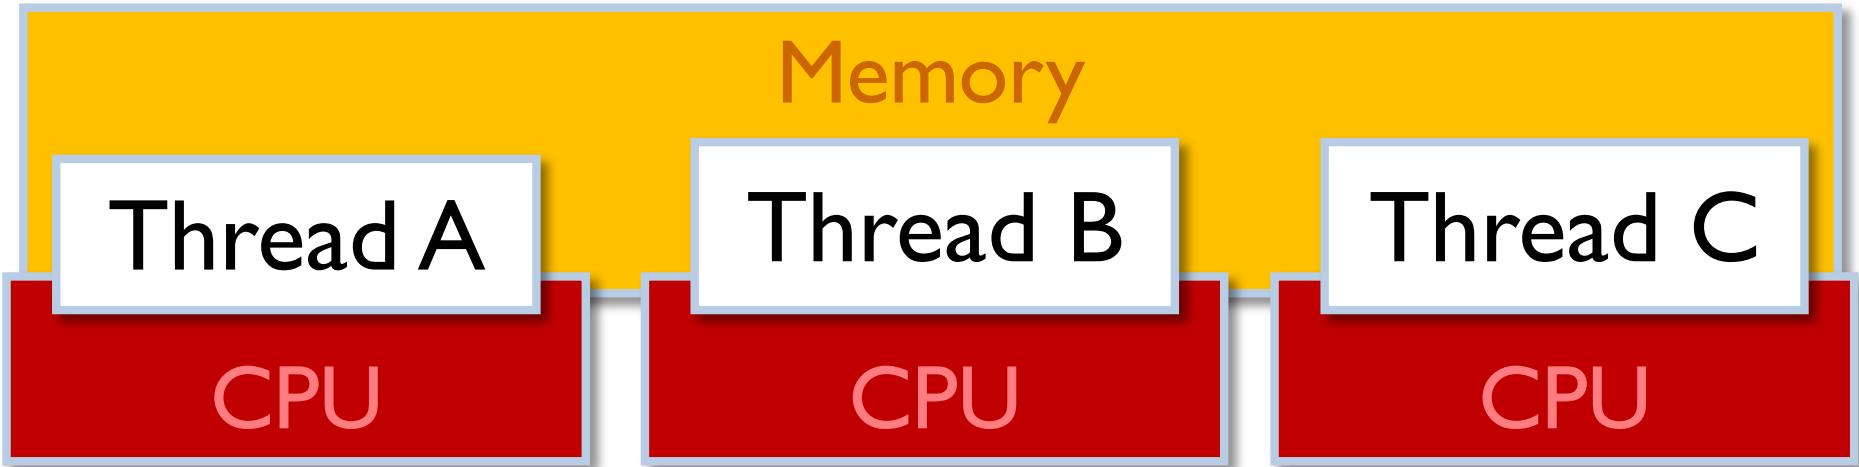
- When a resource is slow, don't want to wait on it
- Windowing system
  - One thread per window, waiting for window input
  - What is slow?
    - Human input, mouse, keyboard
- Network file/web/DB server
  - One thread per incoming request
  - What is slow?
    - Network, disk, remote user (e.g. ATM bank customer)

# Where are threads used?

- When a resource is slow, don't want to wait on it
- Operating system kernel (central component of the OS)
  - One thread waits for keyboard input
  - One thread waits for mouse input
  - One thread writes to the display
  - One thread writes to the printer
  - One thread receives data from the network card
  - One thread per disk ...
  - *Just about everything except the CPU is slow*

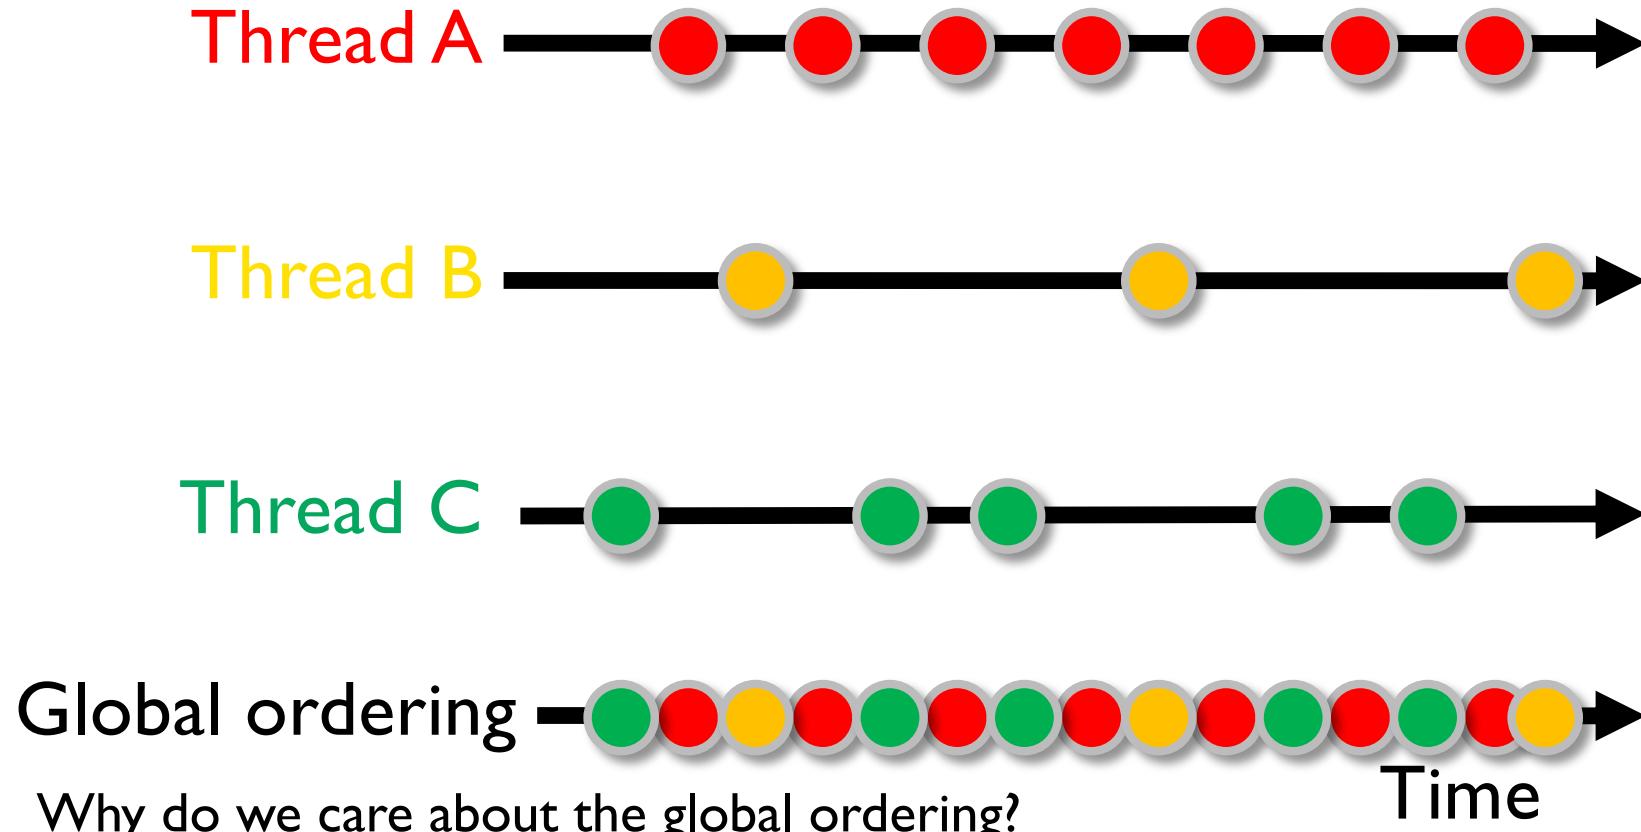
# Cooperating threads

- Assume each thread has its own CPU
  - We will relax this assumption later



- CPUs run at unpredictable speeds
  - Source of non-determinism

# Non-determinism and Ordering



Why do we care about the global ordering?

- Might have dependencies between events

- Different orderings can produce different results

- Not all operations “commutative” (actually very few)

Why is this ordering unpredictable?

- Can't predict how fast processors will run

# Non-determinism example I

- **Thread A:** cout << “ABC”;
- **Thread B:** cout << “123”;
- Possible outputs?
  - “AIBC23”, “ABC123”, ...
- Impossible outputs? Why?
  - “321CBA”, “B12C3A”, ...
- What is shared between threads?
  - Screen, maybe the output buffer

# Non-determinism example 2

- $y = 10;$
- **Thread A:** `int x = y+1;`
- **Thread B:** `y = y*2;`
- Possible results?
  - **A** goes first:  $x = 11$  and  $y = 20$
  - **B** goes first:  $y = 20$  and  $x = 21$
- What is shared between threads?
  - Variable  $y$

# Non-determinism example 3

- `x = 0;`
- **Thread A:** `x = 1;`
- **Thread B:** `x = 2;`
- Possible results?
  - **B** goes first: `x = 1`
  - **A** goes first: `x = 2`
- Is `x = 3` possible?

# Example 3, continued

- What if “ $x = <\text{int}>;$ ” is implemented as
  - $x := x \& 0$
  - $x := x | <\text{int}>$
- Consider this schedule
  - Thread A:  $x := x \& 0$
  - Thread B:  $x := x \& 0$
  - Thread B:  $x := x | 1$
  - Thread A:  $x := x | 2$       **Result:  $x=3!$**
- The “equals” operator is actually two steps!

# Atomic operations

- Must know what operations are **atomic**
  - ...before we can reason about cooperation
- **Atomic**
  - Indivisible
  - Happens without interruption
- Between start and end of atomic action...
  - ...no events from other threads can occur

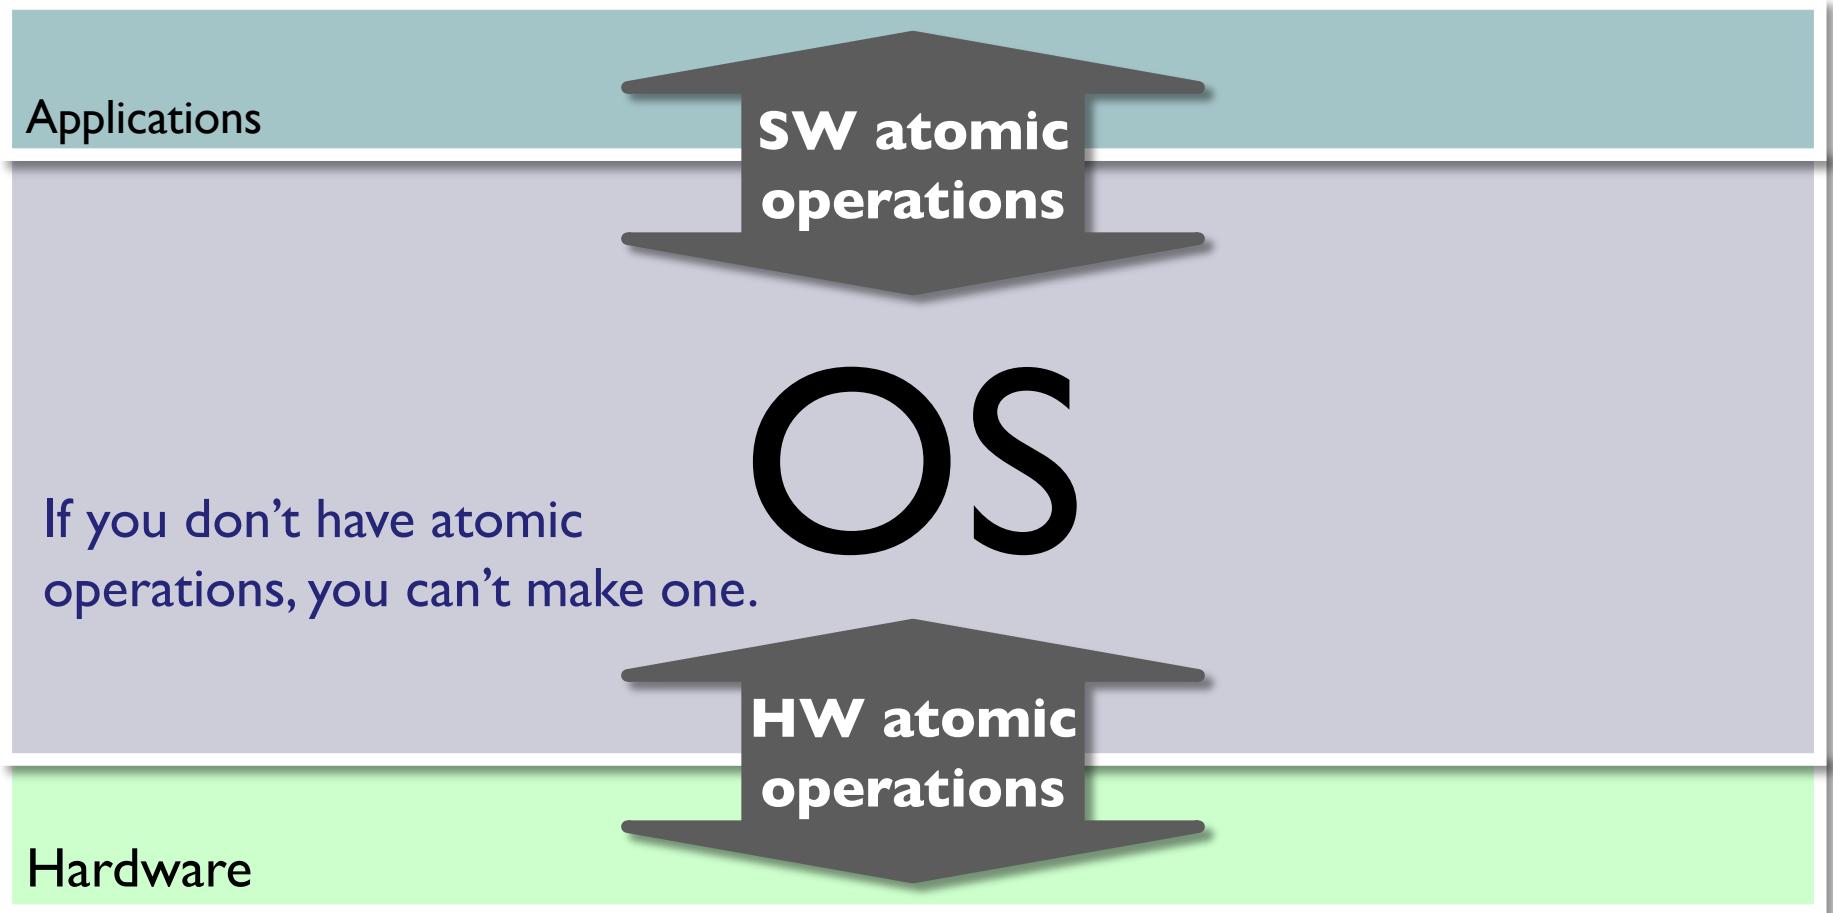
# Review of examples

- Print example (ABC, 123)
  - What did we assume was atomic?
  - What if “print” is atomic?
  - What if printing a char was not atomic?
- Arithmetic example ( $x=y+1$ ,  $y=y*2$ )
  - What did we assume was atomic?

# Atomicity in practice

- On *most* (but not all) machines
  - Memory assignment/reference is atomic
  - E.g.:  $a=1$ ,  $a=b$
- Many other instructions are not atomic
  - E.g.: double-precision floating point store
  - (often involves two memory operations)

# Virtual/physical interfaces



# Another example

- Two threads (A and B)
  - A tries to increment i
  - B tries to decrement i

Thread A:

```
i = 0;  
while (i < 10){  
    i++;  
}  
print "A done."
```

Thread B:

```
i = 0;  
while (i > -10){  
    i--;  
}  
print "B done."
```

# Example continued

- Who wins?
- Does someone have to win?

Thread A:

```
i = 0;  
while (i < 10){  
    i++;  
}  
print "A done."
```

Thread B:

```
i = 0;  
while (i > -10){  
    i--;  
}  
print "B done."
```

# Example continued

- Will it go on forever if both threads:
  - Start at about the same time
  - And execute at exactly the same speed?
  - Yes, if each C *statement* is atomic.

Thread A:

```
i = 0;  
while (i < 10){  
    i++;  
}  
print "A done."
```

Thread B:

```
i = 0;  
while (i > -10){  
    i--;  
}  
print "B done."
```

# Example continued

- What if `i++/i--` are not atomic?
  - `tmp := i+1`
  - `i := tmp`
  - (`tmp` is private to A and B)

# Example continued

- Non-atomic `i++/i--`
  - If A starts  $\frac{1}{2}$  statement ahead, B can win
- How?

Thread A: `tmpA := i + 1 // tmpA == 1`

Thread B: `tmpB := i - 1 // tmpB == -1`

Thread A: `i := tmpA // i == 1`

Thread B: `i := tmpB // i == -1`

# Example continued

- Non-atomic `i++/i--`
  - If A starts  $\frac{1}{2}$  statement ahead, B can win
- How?
- Do you need to worry about this?
  - Yes!!! No matter how unlikely

# Debugging non-determinism

- Requires **worst-case** reasoning
  - Eliminate **all** ways for program to break
- Debugging is hard
  - Can't test all possible interleavings
  - Bugs may only happen sometimes
- **Heisenbug**
  - Re-running program may make the bug disappear
  - Doesn't mean it isn't still there!