

# ECE 570/670

David Irwin  
Lecture 4

# Administrative Details

- Web: <https://courses.umass.edu/eceng570-deirwin/spring23>
  - Username: **irwin-ece**, Password: **myclass**
- Assignment 0 is now posted (see calendar for due date)
  - Look at Assignment 0, find assignment partner (if desired), request username/password, start assignment
  - Will register new groups with autograder over weekend
- UNIX paper assigned for next class

# ECE670 Paper Reviews

- Paper reviews starting now
  - Should do 10 of 12 for class
- Sign up for account linked of Assignments page
  - Only 7 currently signed up
  - Everyone must sign up before I can assign reviews!
- Paper Review
  - Briefly summarize Paper
  - What are the paper's strengths and contributions?
  - Why do you think this paper is considered novel/important?
    - Most of these papers are seminal, so they have few weaknesses
    - They were already accepted at the top venues
- See: <https://people.inf.ethz.ch/troscoe/pubs/review-writing.pdf>

# Assignment 0

- **Inverted Index**

- Virtual appliance posted (see Resources page)
- C/C++ references posted (also see Resources page)
  - ...and Google
- Autograder information posted
  - See link at bottom of Assignment 0 page

- **Configuration**

- Need to know assignment partners ASAP
  - Will use them to configure the autograder
- 7 people have yet to sign up

# Assignment 0

- **Inverted Index**

- **Input:** one argument → a file name
  - Contains list of other text file names
  - Each text file has words, i.e., ASCII characters, in it
- **Output:** all words in alphabetical order, followed by document numbers in which they appear

- **Example**

- `./inverter index.txt`
- `index.txt` contains `foo1.txt` and `foo2.txt`
  - `foo1.txt` contains “this is a test. cool.”
  - `foo2.txt` contains “this is also a test.\nboring.”

# Assignment 0

- **Inverted Output**

a: 0 |

also: 1

boring: 1

cool: 0

foo1.txt → Document 0

is: 0 |

foo2.txt → Document 1

test: 0 |

this: 0 |

- **To compile:** \$> g++ inverter.cc –o inverter

- **To run:** \$> ./inverter

- Debug using gdb – see web for tutorials

- Free to develop anywhere, but **must run and submit** from virtual appliance

# Assignment 0

- **Virtual Appliance**

- Download ECE670-Spring23.ova file from Resources page
  - Requires 2GB of space; will send out passwords via email
- Install VirtualBox (works with Windows, Mac, Linux)
- Click “File”; then “Import Appliance”; click \*.ova file
- Start VM and login with supplied username/password

- **Autograder**

- ***The autograder is merciless! Start early!***
- 1 submission per day plus 3 bonus submissions
- Receive email of results to @umass addr within 4 hours
- To submit: **\$> submit570 0 inverter.cc**
  - Make sure you are connected to the Internet first

# Assignment 0

- **Reminder: do not put code on public online repository, e.g., GitHub, BitBucket, Google Code, etc.!**
- **Assignment 0 is meant to be easy**
  - Do not be lulled into a false sense of security
  - Assignments get progressively more challenging
  - Purpose of this assignment:
    - 1. Understand how the autograding system works
    - 2. Get a refresher on coding in C/C++

# Assignment 0

- **Pro Tip #1**

- Use “scp” to copy files to/from your VM
  - scp == **S**ecure **C**opy
- **In Mac:**
  - 1. Open Applications->Utilities->Terminal
  - 2. scp -P 2500 file.txt group<#>@localhost:.
  - 3. scp -P 2500 group<#>@localhost:file.txt .
  - Will need to enter username and password
- **In Windows:**
  - Use WinSCP program posted on website
  - Use Windows Shell

# Assignment 0

- **Pro Tip #1 (cont'd)**

- “Port forwarding” via port 2500 should already be setup
  - This is why you put the “-P 2500” arguments after scp
- If scp doesn’t work, and port forwarding is not setup, there are directions posted on the Resources part of the web site to configure it

# Assignment 0

## •**Pro Tip #2**

- If your VM is too slow, don't use the GUI
- **ssh** into it from your laptop instead
  - ssh = **S**ecure **S**hell
- **In Mac:**
  - Open Applications->Utilities->Terminal
  - Run \$>ssh -p 2500 group<#>@localhost
  - Enter username and password
- **In Windows:**
  - Use Putty program posted on web or...
  - Windows Shell

# 670 Projects

- **Proposals due Friday March 24th**
  - Requirements on website (2 page proposal)
    - Must have a **software implementation** component
    - Must have a **performance evaluation** component
  - Form into groups (up to 4) and send me names
    - Will post a calendar for group meetings
- **Stock project:** design a web server
  - Implement threaded and/or event-driven web server
  - Evaluate and compare their performance

# Last Time

- Discussed events and threads

# Why Events Are A Bad Idea (for high-concurrency servers)

- Conventional Wisdom
  - Threads “easy” to program....
  - ...but don’t perform well at scale
    - E.g., with *100,000 threads*
- Why?
  - Performance, Control flow, Synchronization, State Management, Scheduling

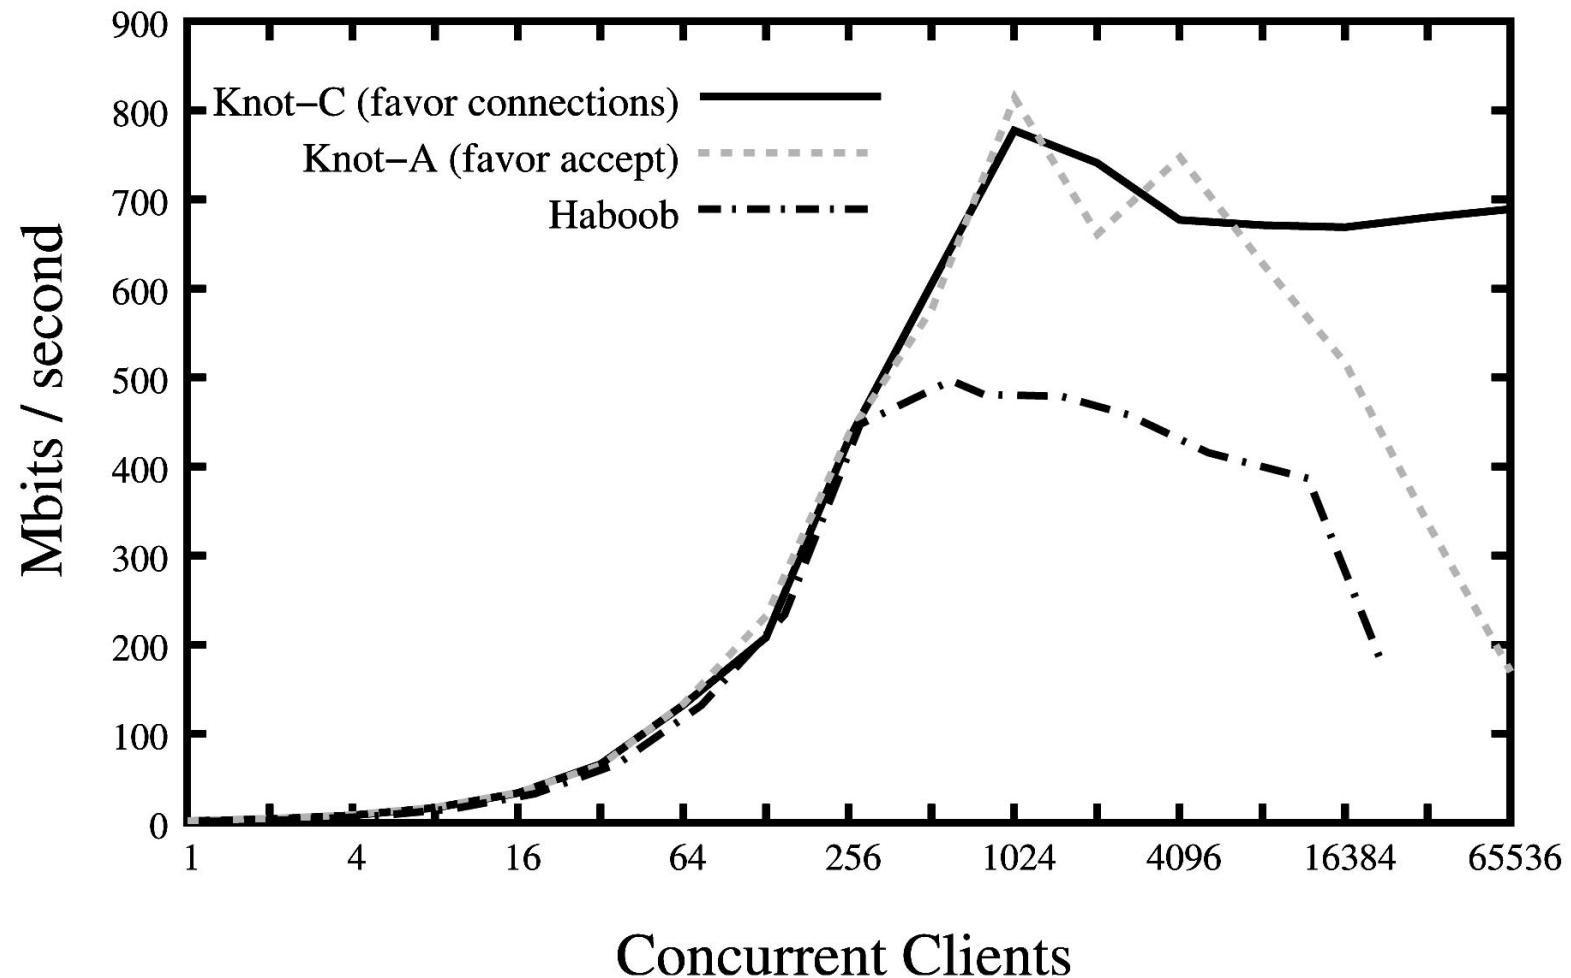
# Why Events Are A Bad Idea (for high-concurrency servers)

- Performance – remove  $O(n)$  operations
- Control Flow – no need for complex non-linear flows
- Synchronization – not easy for either on multiprocessors
- State Management – stacks easier to manage
- Scheduling – scheduling tricks applicable to threads

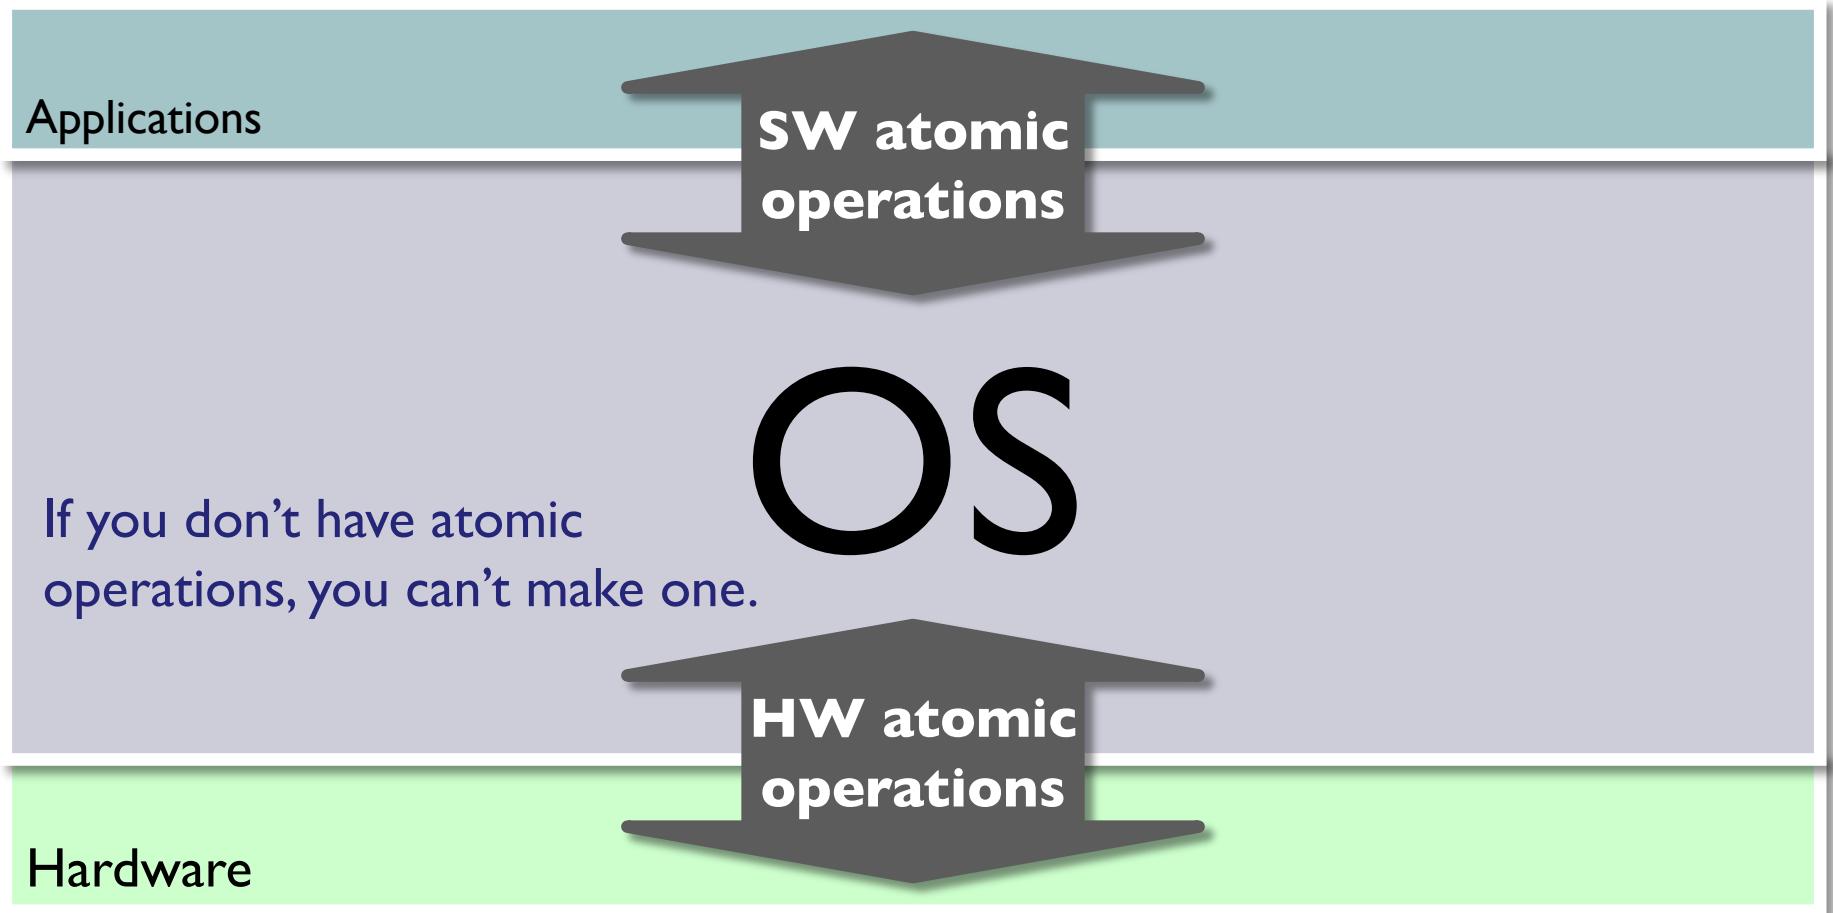
# Why Events Are A Bad Idea (for high-concurrency servers)

- Directions:
  - Dynamic Stack Growth (change stack size dynamically)
  - Live State Management (automatically purge stack of unnecessary state)
  - Synchronization tools (use automated tools to find synchronization bugs)

# Why Events Are A Bad Idea (for high-concurrency servers)



# Virtual/physical interfaces



# Debugging non-determinism

- Requires **worst-case** reasoning
  - Eliminate **all** ways for program to break
- Debugging is hard
  - Can't test all possible interleavings
  - Bugs may only happen sometimes
- **Heisenbug**
  - Re-running program may make the bug disappear
  - Doesn't mean it isn't still there!

# Constraining concurrency

- **Synchronization**
  - Controlling thread interleavings
- Some events are independent
  - Don't depend on shared state
  - Relative order of these events don't matter
- Other events are dependent
  - Output of one can be input to another
  - Their order can affect program results

# Goals of synchronization

## I. All interleavings must give correct result

- Correct concurrent program
  - Works no matter how fast threads run
  - Important for your projects!

## 2. Constrains program as little as possible

- Why?
  - Constraints slow program down
  - Constraints create complexity

# “Too much milk” principal



# “Too much milk” rules

- The fridge must **always** be stocked with milk
  - Milk expires quickly, so never >1 milk. We need exactly 1.
- David and Jeannie (two threads)
  - Can come home at any time
  - If either sees an empty fridge, must immediately go to store and buy milk
  - Code below (no synchronization)

```
if (noMilk){  
    buy milk;  
}
```



# Unsynchronized code will break

Time		
3:00	Look in fridge (no milk)	
3:05	Go to grocery store	
3:10		Look in fridge (no milk)
3:15	Buy milk	
3:20		Go to grocery store
3:25	Arrive home, stock fridge	
3:30		Buy milk
3:35		Arrive home, stock fridge Too much milk!



# What broke?

- Code worked sometimes, but not always
  - Code contained a **race condition**
  - Processor speed caused incorrect result
- First type of synchronization
  - **Mutual exclusion**
  - **Critical sections**

# Synchronization concepts

- **Mutual exclusion**
  - Ensure only one thread doing something at a time
    - *i.e., one person shops at a time*
  - Code blocks atomic with respect to each other
  - Threads can't run code blocks at same time

# Synchronization concepts

- **Critical section**
  - Code “block” that must run atomically
    - “with respect to some other pieces of code”
  - If A and B are critical with respect to each other
    - Threads must not interleave events in A and B
    - A and B *mutually exclude* each other
  - Often conflicting code is in the same block
    - But executed by different threads
    - Reads/writes of shared data (e.g. screen, fridge)

# Back to “Too much milk”

- What is the critical section?

```
if (noMilk){  
    buy milk;  
}
```

- Jeannie and David’s critical sections
  - Must be atomic with respect to each other

# “Too much milk” Solution I

- Assume only atomic load/store
  - Build larger atomic section from load/store
- **Idea:**
  1. Leave notes to say you’re taking care of it
  2. Don’t check milk if there is a note

# Solution I code

- Atomic operations
  - *Atomic load*: check note
  - *Atomic store*: leave note

```
if (noMilk) {  
    if (noNote){  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

# Does it work?



```
if (noMilk) {  
    1 if (noNote){  
        leave note;  
        3 buy milk;  
        remove note;  
    }  
}  
}
```

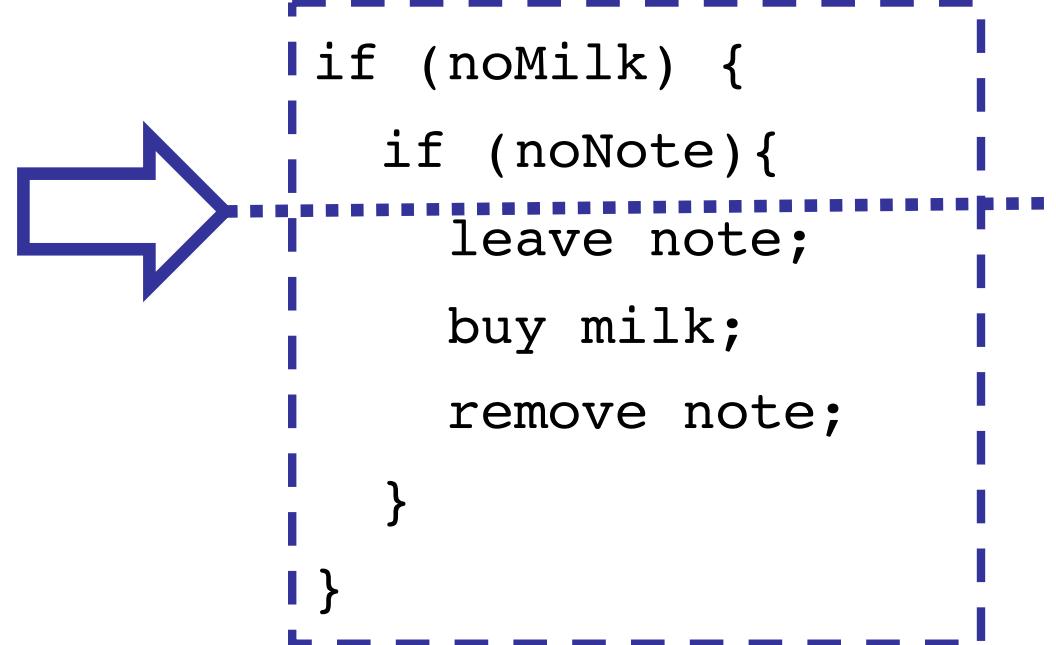
```
if (noMilk) {  
    2 if (noNote){  
        leave note;  
        4 buy milk;  
        remove note;  
    }  
    }  
}
```

Is this better than no synchronization at all?

What if “if” sections are switched?

# What broke?

- David's events can happen
  - After Jeannie checks for a note, but...
  - before Jeannie leaves a note



# Next solution

- **Idea:**
  - Change the order of “leave note”, “check note”
  - Requires labeled notes
  - Why?
    - (You’ll see your note and never do anything)

# Does it work?



```
leave noteJeannie  
if (no noteDavid){  
    if (noMilk){  
        buy milk;  
    }  
}  
remove noteJeannie
```

```
leave noteDavid  
if (no noteJeannie){  
    if (noMilk){  
        buy milk;  
    }  
}  
remove noteDavid
```

Nope. Why?  
Illustration of “starvation.”

# What about now?



```
while (noMilk){  
    leave noteJeannie  
    if(no noteDavid){  
        if(noMilk){  
            buy milk;  
        }  
    }  
    remove noteJeannie  
}
```

```
while (noMilk){  
    leave noteDavid  
    if(no noteJeannie){  
        if(noMilk){  
            buy milk;  
        }  
    }  
    remove noteDavid  
}
```

Nope. (Same starvation problem as before)

# Next solution

- We're getting closer
- Problem
  - Who buys milk if both leave notes?
- Solution
  - Let Jeannie hang around to make sure job is done

# Does it work?



```
leave noteJeannie  
while (noteDavid){  
    do nothing  
}  
if (noMilk){  
    buy milk;  
}  
remove noteJeannie
```

```
leave noteDavid  
if (no noteJeannie){  
    if (noMilk){  
        buy milk;  
    }  
}  
remove noteDavid
```

Yes! It does work! Can you prove it?

# Downside of solution

- Complexity
  - Hard to convince yourself it works
- Asymmetric
  - Jeannie and David run different code
- Not clear if this scales to > 2 people
- Jeannie consumes CPU while waiting
  - `while (noteDavid) { do nothing }`
  - **Busy-waiting**
- **But:** only needed atomic load/store!

# Raising the level of abstraction

- Mutual exclusion with atomic load/store
  - Painful to program
  - Wastes resources (busy waiting)
    - Need more HW support to fix this
    - Will be covered later
- OS can provide higher level abstractions