

# ECE 570/670

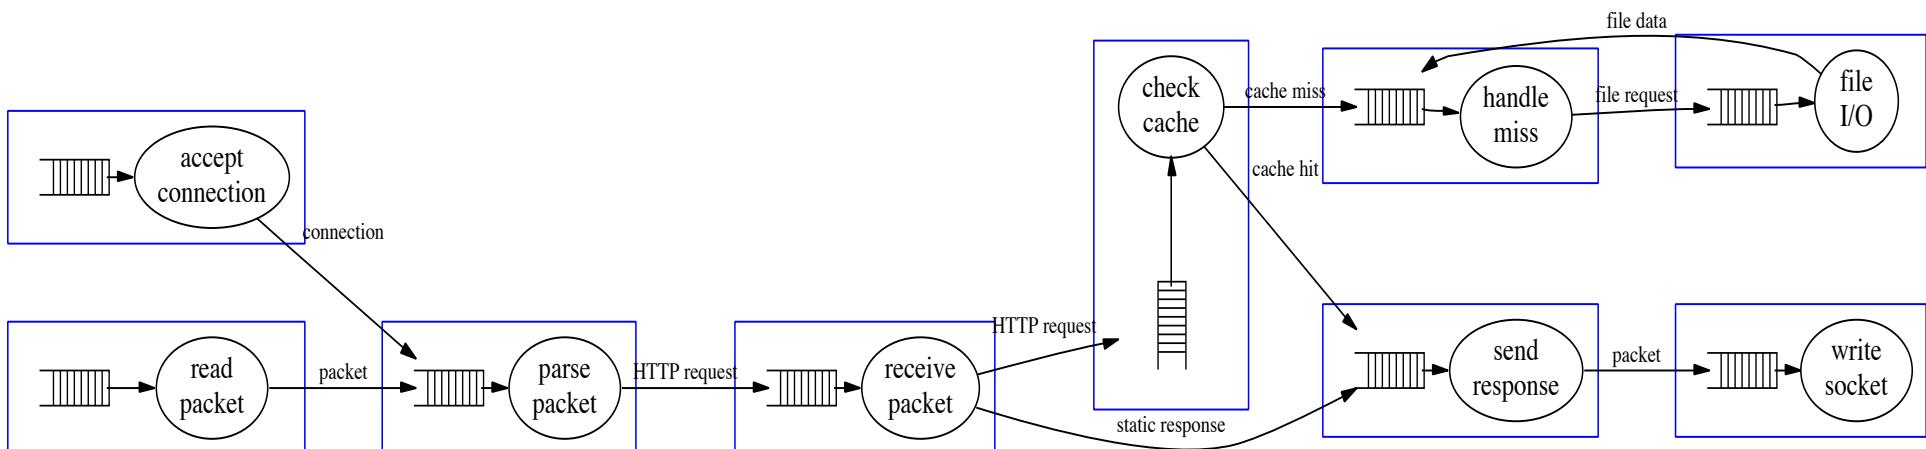
David Irwin  
Lecture 6

# Outline

- Virtualization paper discussion
  - Paper posted for Thursday
  - **End to End Arguments in Systems Design**
- ECE670 project signups
  - Only 1 group signed up so far
- Continue synchronization
- Maybe get to Assignment I discussion

# Virtualization Considered Harmful: OS Design Directions for Well-Conditioned Services

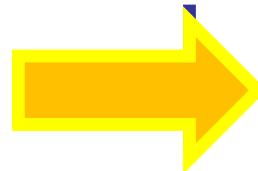
- New server architecture
  - Staged Event Driven Architecture (SEDA)
  - Set of event-driven stages separated by queues
- Benefits
  - Small number of threads (high concurrency)
  - Queues individually “load-conditioned”
  - Modularity of stages (aids debugging)



# How about this?

I'm always  
holding a lock  
while accessing  
shared state.

ptr may not point to  
tail after unlock/lock.



```
enqueue () {  
    lock (qLock)  
    node *ptr;  
    for (ptr=head;  
         ptr->next!=NULL;  
         ptr=ptr->next) {}  
    unlock (qLock);  
    lock (qLock);  
    ptr->next=new_element;  
    new_element->next=NULL;  
    unlock(qLock);  
}
```

## Lesson:

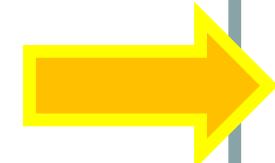
- Thinking about individual accesses is not enough
- Must reason about dependencies *between* accesses too

# Next up

- Mutual exclusion is great and all, but...
  - It's not really expressive enough
- Ordering constraints
  - Often must wait for something to happen
  - Introduce monitors and condition variables

# Intro to ordering constraints

- Say you want dequeue to wait while the queue is empty
- Can we just busy-wait?
  - No!
  - Still holding lock



```
dequeue () {  
    lock (qLock);  
    element=NULL;  
    while (head==NULL) {}  
    // remove head  
    element=head->next;  
    head->next=NULL;  
    unlock (qLock);  
    return element;  
}
```

# Release lock before spinning?

What can go wrong?

- Another dequeuer could “steal” our element
- Head might be NULL when we try to remove entry



```
dequeue () {  
    lock (qLock);  
    element=NULL;  
    unlock (qLock);  
    while (head==NULL) {}  
    lock (qLock);  
    // remove head  
    element=head->next;  
    head->next=NULL;  
    unlock (qLock);  
    return element;  
}
```

# One more try

- Does it work?
  - Seems ok
- Why?
  - Shared state is protected
- What's wrong?
  - Busy-waiting
  - Wasteful

```
dequeue () {  
    lock (qLock);  
    element=NULL;  
    while (head==NULL) {  
        unlock (qLock);  
        lock (qLock);  
    }  
    // remove head  
    element=head->next; ←  
    head->next=NULL;  
    unlock (qLock);  
    return element;  
}
```

# Ideal solution

- Would like dequeuing thread to “sleep”
  - Add self to “waiting list”
  - Enqueuer can wake up waiting threads when queue is non-empty
- **Problem:** what to do with the lock?
  - Can dequeuing thread sleep with lock?
    - No!
    - Enqueuer would never be able to add to queue

# Release the lock before sleep?

```
enqueue () {  
    acquire lock  
    find tail of queue  
    add new element  
    if (dequeuer waiting){  
        remove from wait list  
        wake up dequeuer  
    }  
    release lock  
}
```

```
dequeue () {  
    acquire lock  
    ...  
    if (queue empty) {  
        release lock  
        add self to wait list  
        sleep  
    }  
    ...  
    release lock  
}
```

Does this work?

# Release the lock before sleep?

enqueue () {  
 acquire lock  
 find tail of queue  
 add new element  
 if (dequeuer waiting){  
 remove from wait list  
 wake up dequeuer  
 }  
 release lock  
}

dequeue () {  
 acquire lock  
 ...  
 if (queue empty) {  
 release lock  
 add self to wait list  
 sleep  
 }  
 ...  
 release lock  
}

Thread can sleep forever!

Other problems?

Wait list is shared and unprotected. (bad idea)

# Release the lock before sleep?

```
enqueue () {  
    acquire lock  
    find tail of queue  
    add new element  
    if (dequeuer waiting){  
        remove from wait list  
        wake up dequeuer  
    }  
    release lock  
}
```

```
dequeue () {  
    acquire lock  
    ...  
    if (queue empty) {  
        add self to wait list  
        release lock  
        sleep  
    }  
    ...  
    release lock  
}
```

# Release the lock before sleep?

enqueue () {  
 acquire lock  
 find tail of queue  
 add new element  
 if (dequeuer waiting){  
 remove from wait list  
 wake up dequeuer  
 }  
 release lock  
}

dequeue () {  
 acquire lock  
 ...  
 if (queue empty) {  
 add self to wait list  
 release lock  
 sleep  
 }  
 ...  
 release lock  
}

**Problem:** missed wake-up

*Note: this can be fixed, but it's messy*

# Two types of synchronization

- **As before we need to raise the level of abstraction**

## 1. Mutual exclusion

- One thread doing something at a time
- Use locks

## 2. Ordering constraints

- Describes “before-after” relationships
- One thread waits for another
- Use monitors

# Review

- First we discussed **atomic actions**
  - i.e., load, store at hardware layer
  - Guaranteed to happen without being interrupted
  - Not “strong” enough to provide mutual exclusion
- Then we discussed **locks/mutexes**
  - Provided mutual exclusion
  - Prevented two threads from accessing shared data simultaneously
  - Commonly used for “identical” threads running in parallel
  - Cannot provide ordering
- Now we are going to discuss **monitors**
  - Provides ordering (before-after relationship)

# Monitor synchronization

## I. Mutual exclusion

- One thread doing something at a time
- Use locks

## 2. Ordering constraints

- Describes “before-after” relationships
- One thread waits for another
- **Monitor:** a lock + its **condition variable**

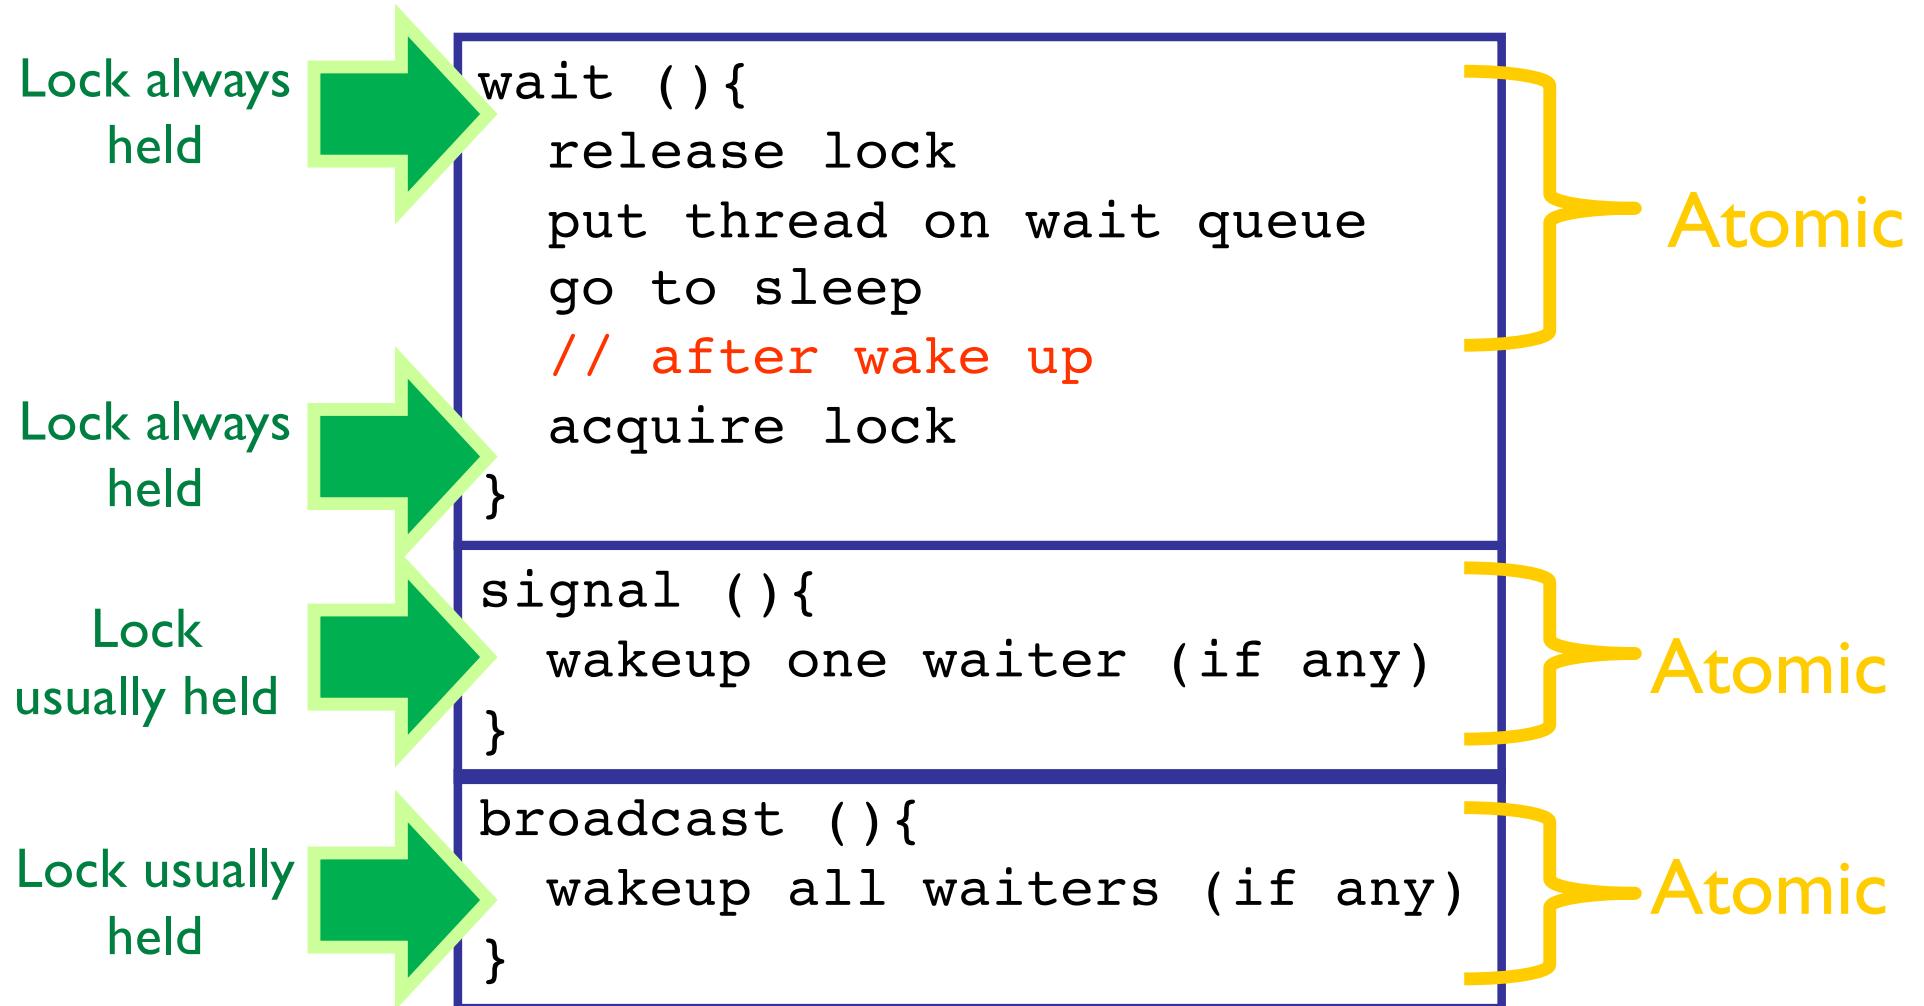
# Locks and condition variables

- **Condition variable**
  - Lets threads sleep inside a critical section
- Internal atomic actions (for now, by definition)

```
// begin atomic
release lock
put thread on wait queue
go to sleep
// end atomic
```

- CV State = queue of waiting threads + one lock

# Condition variable operations



# CVs and invariants

- User programs
  - Ok to leave invariants violated before wait?
  - **No: wait releases the lock**
- Larger rule about returning from wait
  - Lock may have changed hands
  - State can change between wait entry and return
  - Don't make assumptions about shared state

# Multi-threaded queue

```
enqueue () {  
    acquire lock  
  
    find tail of queue  
    add new element  
  
    signal (lock, CV)  
  
    release lock  
}
```

```
dequeue () {  
    acquire lock  
  
    if (queue empty) {  
        wait (lock, CV)  
    }  
  
    remove item from queue  
    release lock  
    return removed item  
}
```

Any problems with the “if” statement in dequeue?

# Multi-threaded queue

```
enqueue () {  
    acquire lock  
  
    find tail of queue  
    add new element  
  
    signal (lock, CV)  
  
    release lock  
}
```

```
dequeue () {  
    acquire lock  
  
    if (queue empty) {  
        // begin atomic wait  
        release lock  
        sleep and wait  
        // end atomic wait  
        re-acquire lock  
    }  
  
    remove item from queue  
    release lock  
    return removed item  
}
```

# Multi-threaded queue

```
enqueue () {  
    acquire lock  
  
    find tail of queue  
    add new element  
  
    signal (lock, CV)  
  
    release lock  
}
```

2

```
dequeue () {  
    acquire lock  
    ...  
    return removed item  
}
```

3

```
dequeue () {  
    acquire lock  
  
    if (queue empty) {  
        // begin atomic wait  
        release lock  
        sleep and wait  
        // end atomic wait  
        re-acquire lock  
    }  
}
```

1

4

```
remove item from queue  
release lock  
return removed item  
}
```

# Multi-threaded queue

```
enqueue () {  
    acquire lock  
  
    find tail of queue  
    add new element  
  
    signal (lock, CV)  
  
    release lock  
}
```

```
dequeue () {  
    acquire lock  
  
    if (queue empty) {  
        // begin atomic wait  
        release lock  
        sleep and wait  
        // end atomic wait  
        re-acquire lock  
    }  
  
    remove item from queue  
    release lock  
    return removed item  
}
```

## How to solve?

# Multi-threaded queue

The “condition” in condition variable

```
enqueue () {  
    acquire lock  
  
    find tail of queue  
    add new element  
  
    signal (lock, CV)  
  
    release lock  
}
```

```
dequeue () {  
    acquire lock  
  
    while (queue empty) {  
        wait (lock, CV)  
    }  
  
    remove item from queue  
    release lock  
    return removed item  
}
```

Solve with a while loop (“loop before you leap”)  
Remember this for assignments ...

# Mesa vs. Hoare monitors

- So far, we've described Mesa monitors
  - After waking up, no special priority
  - Threads have to recheck condition
- Hoare semantics are “simpler”
  - But complicate implementation

# Hoare semantics

- Condition guaranteed true after wakeup
  - i.e. no need to loop
- Waiter acquires lock before any threads run
  - (since lock protects condition)
  - Including the signaler!
- Signaler must give up lock and signal atomically
- What would this mean for invariants?
  - All invariants must be established before signal
- **We will use Mesa semantics**

# Tips for using monitors

- 1. List the shared data needed for problem**
- 2. Figure out the locks**
  - 1 lock per group of shared data
- 3. Bracket code that uses shared data with lock/unlock**
- 4. Think about before-after conditions**
  - 1 condition variable per type of condition
  - CV's lock should protect data used to evaluate condition
  - Call wait when you need to wait on a condition
  - Loop to re-check condition when wait returns
  - Call signal when condition changes
  - Ensure invariants are established when lock is not held (unlock, wait)

# Monitor Benefits

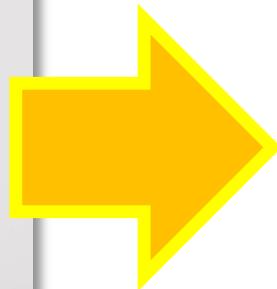
- Benefit of sleeping thread vs. blocked thread?
  - Sleeping thread is not wasting CPU cycles
- A note about condition variables
  - They represent some condition in our code
  - But they are not boolean expressions themselves
  - When we call wait or signal, we specify mutex and condition variable

# Producer-consumer problem

- Producer makes something consumer wants
  - Goal: avoid lock-step (direct hand-off)



Delivery person  
(producer)



Soda drinker  
(consumer)

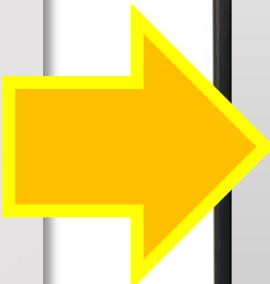


Problem: everyone's time is wasted.

# Producer-consumer problem

- Instead, use a fixed-size, shared buffer (sometimes called bounded buffer)
  - Producer puts in (must wait when full)
  - Consumer takes out (must wait when empty)
  - Must synchronize access to buffer
- Examples
  - Unix pipes: cat file | grep -v the | grep from | wc -l
  - Interaction between hardware devices/buffers

# Use a soda machine as a buffer



# Solving producer-consumer

## 1. What are the variables/shared state?

- Soda machine buffer
- Number of sodas in machine ( $\leq$  MaxSodas)

## 2. Locks?

- 1 to protect all shared state (sodaLock)

## 3. Mutual exclusion?

- Only one thread can manipulate machine at a time

## 4. Ordering constraints?

- Consumer must wait if machine is empty (CV hasSoda)
- Producer must wait if machine is full (CV hasRoom)

# Producer-consumer code

```
consumer () {  
    lock (sodaLock)  
  
    while (numSodas == 0) {  
        wait (sodaLock, hasSoda)  
    }  
  
    take soda out of machine  
  
    signal (hasRoom)  
  
    unlock (sodaLock)  
}  
  
producer () {  
    lock (sodaLock)  
  
    while(numSodas==MaxSodas){  
        wait (sodaLock, hasRoom)  
    }  
  
    add soda to machine  
  
    signal (hasSoda)  
  
    unlock (sodaLock)  
}
```

# Variations: looping producer

- Producer
  - Infinite loop ok?
  - Why/why not?
- Release lock in  
wait call

```
producer () {  
    lock (sodaLock)  
    while (1) {  
        while (numSodas==MaxSodas) {  
            wait (sodaLock, hasRoom)  
        }  
  
        add soda to machine  
  
        signal (hasSoda)  
    }  
    unlock (sodaLock)  
}
```

# Variations: resting producer

- Producer
  - Sleep ok?
  - Why/why not?
- Shouldn't hold locks during a slow operation

```
producer () {  
    lock (sodaLock)  
    while (1) {  
        sleep (1 hour)  
        while (numSodas==MaxSodas) {  
            wait (sodaLock, hasRoom)  
        }  
  
        add soda to machine  
  
        signal (hasSoda)  
    }  
    unlock (sodaLock)  
}
```

# Variations: one CV?

```
consumer () {  
    lock (sodaLock)  
  
    while (numSodas == 0) {  
        wait (sodaLock,hasRorS)  
    }  
  
    take soda out of machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

```
producer () {  
    lock (sodaLock)  
  
    while(numSodas==MaxSodas){  
        wait (sodaLock,hasRorS)  
    }  
  
    add soda to machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

Two producers, two consumers: who consumes a signal?

# Variations: one CV?

```
consumer () {  
    lock (sodaLock)  
  
    while (numSodas == 0) {  
        wait (sodaLock,hasRorS)  
    }  
  
    take soda out of machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

```
producer () {  
    lock (sodaLock)  
  
    while(numSodas==MaxSodas){  
        wait (sodaLock,hasRorS)  
    }  
  
    add soda to machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

Is it possible to have a producer and consumer both waiting?  
max=1, cA and cB wait, pC adds/signals, pD waits, cA wakes

# Variations: one CV?

```
consumer () {  
    lock (sodaLock)  
  
    while (numSodas == 0) {  
        wait (sodaLock,hasRorS)  
    }  
  
    take soda out of machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

```
producer () {  
    lock (sodaLock)  
  
    while(numSodas==MaxSodas){  
        wait (sodaLock,hasRorS)  
    }  
  
    add soda to machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

Is it possible to have a producer and consumer both waiting?  
max=1, cA and cB wait, pC adds/signals, pD waits, cA wakes  
cA removes/signals, cB wakes...but numSodas==0

# Variations: one CV?

```
consumer () {  
    lock (sodaLock)  
  
    while (numSodas == 0) {  
        wait (sodaLock,hasRorS)  
    }  
  
    take soda out of machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

```
producer () {  
    lock (sodaLock)  
  
    while(numSodas==MaxSodas){  
        wait (sodaLock,hasRorS)  
    }  
  
    add soda to machine  
  
    signal (hasRorS)  
  
    unlock (sodaLock)  
}
```

How can we make the one CV solution work?

# Variations: one CV?

```
consumer () {  
    lock (sodaLock)  
  
    while (numSodas == 0) {  
        wait (sodaLock,hasRorS)  
    }  
  
    take soda out of machine  
  
    broadcast (hasRorS)  
  
    unlock (sodaLock)  
}
```

```
producer () {  
    lock (sodaLock)  
  
    while(numSodas==MaxSodas){  
        wait (sodaLock,hasRorS)  
    }  
  
    add soda to machine  
  
    broadcast (hasRorS)  
  
    unlock (sodaLock)  
}
```

Use broadcast instead of signal

# Broadcast vs signal

- Can I always use broadcast instead of signal?
  - Yes, assuming threads recheck condition
- Why might I use signal instead?
  - Efficiency (spurious wakeups)
  - May wakeup threads for no good reason

# Java Monitors

```
class BoundedBuffer {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull =  
        lock.newCondition();  
    final Condition notEmpty =  
        lock.newCondition();  
  
    final Object[] items = new Object[100];  
    int putptr, takeptr, count;  
  
    public void put(Object x)  
        throws InterruptedException {  
        lock.lock();  
        try {  
            while (count == items.length)  
                notFull.await();  
            items[putptr] = x;  
            if (++putptr == items.length)  
                putptr = 0;  
            count++;  
            notEmpty.signal();  
        } finally {  
            lock.unlock();  
        }  
    }  
  
    public Object take()  
        throws InterruptedException {  
        lock.lock();  
        try {  
            while (count == 0)  
                notEmpty.await();  
            Object x = items[takeptr];  
            if (++takeptr == items.length)  
                takeptr = 0;  
            count--;  
            notFull.signal();  
            return x;  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

# C/C++ Monitors

```
#include <pthread.h>

pthread_mutex_t mutex =
    PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t notFull =
    PTHREAD_COND_INITIALIZER;
pthread_cond_t notEmpty =
    PTHREAD_COND_INITIALIZER;

int items[100];
int putptr, takeptr, count = 0;

void put(int x) {
    pthread_mutex_lock( &mutex );
    while( count == 100 )
        pthread_cond_wait( &notFull, &mutex );
    items[putptr] = x;
    if (++putptr == 100) putptr = 0;
    count++;
    pthread_cond_signal( &notEmpty );
    pthread_mutex_unlock( &mutex );
}

int take() {
    pthread_mutex_lock( &mutex );
    while (count == 0)
        pthread_cond_wait( &notEmpty, &mutex );
    int x = items[takeptr];
    if (++takeptr == 100) takeptr = 0;
    count--;
    pthread_cond_signal( &notFull );
    pthread_mutex_unlock( &mutex );
    return x;
}

int main(int argc, char* argv[ ]) {
    return 0;
}
```

g++ buffer.cc -pthread