

ECE 570/670

David Irwin
Lecture 7

Administrative

- Assignment 1 out
- Assignment 0 due
 - 2 people have yet to sign up with the autograder
- ECE670 Project groups
 - 16 have yet to fill out Google Form
 - Will send reminder email today
 - Will send out appointment calendar soon

Administrative

- New TA
 - Xiaoding (Rebecca) Guan
 - Will post office hours and info soon

Administrative

- Today
 - End-to-End argument
 - Assignment I
 - Higher-level synchronization

End-to-End Argument

- *In a general-purpose network, application-specific functions ought to reside in end hosts rather than intermediary nodes, provided that they can be implemented “completely and correctly” in the end hosts.*
- **Benefit of adding functions diminishes quickly, especially when hosts must re-implement functions for completeness/correctness**
- Since any function incurs some penalty even when not used, implementing it in the network distributes the penalty across all clients, regardless of whether they use it

End-to-End Argument

- *In a general-purpose network, application-specific functions ought to reside in end hosts rather than intermediary nodes, provided that they can be implemented “completely and correctly” in the end hosts.*
- Benefit of adding functions diminishes quickly, especially when hosts must re-implement functions for completeness/correctness
- **Since any function incurs some penalty even when not used, implementing it in the network distributes the penalty across all clients, regardless of whether they use it**

End-to-End Argument

- **Example: File transfer program**
 - Providing reliable data communication does not reduce the burden of the application to perform checksum/retries
 - MIT example: made too strong assumptions about network reliability
 - Performance Optimization
 - End-to-end recovery cost increases with file size
 - Could increase network reliability for performance
 - Tradeoff b/t overhead and performance benefit
 - Will still need end-to-end check
 - Considerations
 - Low-level optimization applies to all applications
 - E.g., might not work well for real-time communication
 - Lower levels have less information than higher levels

End-to-End Argument

- **Other End-to-End Examples**
 - Delivery guarantees - end-to-end acknowledgements
 - Secure data transmission – end-to-end encryption
 - Duplicate message suppression – application may generate duplicates that the network cannot distinguish; requires end-to-end suppression
 - FIFO message delivery guarantee – messages along different paths cannot guarantee FIFO ordering; have to re-order end-to-end
- **Widely applicable to other areas** – mentions library OSes at the end

Assignment I

- Assignment I is out: two parts
 - 1. Assignment I disk scheduler (Id) (due 3/21)
 - How to use concurrency primitives
 - 2. Assignment I thread library (It) (due later)
 - How concurrency primitives actually work by actually building them
 - Both parts treated as separate assignments with equal weight in final grade
- Start early; you will need the entire time

Assignment I Overview

- Thread library interface
 - `thread_libinit()`
 - `thread_create()`
 - `thread_yield()`
 - `thread_lock()`
 - `thread_unlock()`
 - `thread_wait()`
 - `thread_signal()`
 - `thread_broadcast()`
- Defined in `thread.h`
 - See description for specific inputs/return values

Assignment I (Part I)

- Given a working thread library (thread.o)
- Disk scheduler
 - Schedules disk I/Os for multiple threads
 - Threads issue I/Os, which queue up at scheduler
 - Queue has a fixed size; must wait if queue full
 - Many threads issue requests; one thread services them
 - Similar to producer/consumer
 - Disk queue not FIFO – instead handles requests in SSTF (Shortest Seek Time First Order)
 - Each request has track number
 - Disk starts at track 0
 - Next request serviced is one that is closest to the current request

Assignment I (Part I)

- Example: “./disk <queue_size> <list of input files>”
 - Assume one input file per thread
 - “./disk 10 disk.in0 disk.in1 disk.in2 disk.in3 disk.in4”

disk.in0	disk.in1	disk.in2	disk.in3	disk.in4
-----	-----	-----	-----	-----
53	914	827	302	631
785	350	567	230	11

Assignment I (Part I)

- After issuing request, requester thread prints:
 - requester 0 track 53
- After servicing request, servicing thread prints
 - service requester 0 track 53

```
requester 0 track 53
requester 1 track 914
requester 2 track 827
service requester 0 track 53
requester 3 track 302
service requester 3 track 302
requester 4 track 631
service requester 4 track 631
requester 0 track 785
service requester 0 track 785
requester 3 track 230
service requester 2 track 827
requester 4 track 11
service requester 1 track 914
requester 2 track 567
service requester 2 track 567
requester 1 track 350
service requester 1 track 350
service requester 3 track 230
service requester 4 track 11
Thread library exiting.
```

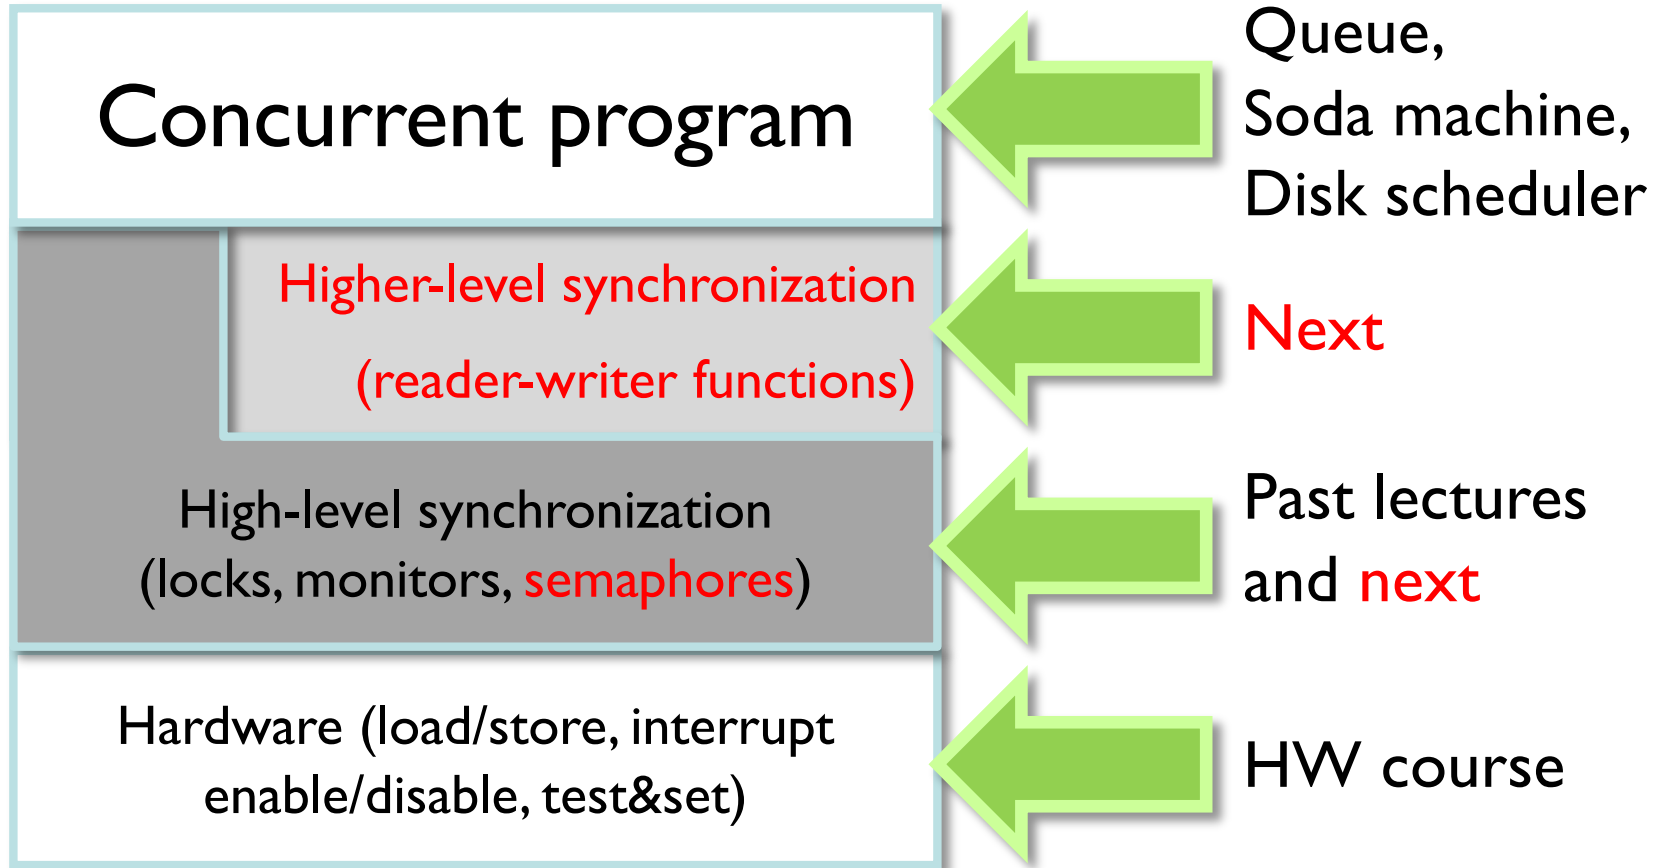
Assignment I (Part I)

- Compiling your program
 - **g++ disk.cc thread.o libinterrupt.a -ldl -o disk.**
- Thread library includes the ability to:
 - Turn preemptions on and off (start_preemptions())
 - Preemptions may be:
 - Asynchronous (happen at random times every 10ms)
 - Synchronous (happen randomly but deterministically on every run)
 - Change change pattern by changing random seeds
 - **These should be used for testing**

Assignment I (Part 2)

- You write the thread library yourself
 - I.e., Implement all of the thread library functions
 - Will use your correct disk scheduler for testing
 - Will discuss more next week
- Get started on the disk scheduler now!
 - It will be challenging and take time
 - **Note: the autograder will not tell you why your program is failing**
 - It is your job to create tests for your program that test various conditions (e.g., boundary values)

Layers of synchronization



Locks thus far

- Lock anytime shared data is read/written
 - Ensures correctness
 - Only one thread can read/write at a time
- Would like more concurrency
- How, without exposing violated invariants?
 - Allow multiple threads to read
 - (as long as none are writing)

Problem definition

- Shared state that many threads want to access
 - Many more reads than writes
 - Peoplesoft, Ebay, any large database
- 1. When no writers, allow multiple readers
- 2. One writer at a time, when no readers
- Goal: build reader-writer locks using monitors

Reader-writer interface

- `readerStart` (called when thread begins reading)
- `readerFinish` (called when thread is finished reading)
- `writerStart` (called when thread begins writing)
- `writerFinish` (called when thread is finished writing)
- If no threads between `writerStart/writerFinish`
 - Many threads between `readerStart/readerFinish`
- Only 1 thread between `writerStart/writerFinish`

Solving reader-writer locks

1. **What are the variables/shared state?**

- Number of readers (numReaders)
- Number of writers (numWriters)

2. **Locks?**

- 1 to protect all shared state (rwLock)

3. **Mutual exclusion?**

- Only one writing thread at a time

4. **Ordering constraints?**

- readerStart must wait if there are writers
- writerStart must wait if there are readers or writers
- Due to overlap, use one CV (ReaderOrWriterFinishCV)

Reader-writer code

```
writerStart () {  
    lock (RWLock)  
    while (numWriters > 0 ||  
           numReaders > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numWriters++;  
    unlock (RWLock)  
}
```

```
writerFinish () {  
    lock (RWLock)  
    numWriters--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

```
readerStart () {  
    lock (RWLock)  
    while (numWriters > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numReaders++  
    unlock (RWLock)  
}
```

```
readerFinish () {  
    lock (RWLock)  
    numReaders--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

Reader-writer code

```
writerStart () {  
    lock (RWLock)  
    while (numWriters > 0 ||  
           numReaders > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numWriters++;  
    unlock (RWLock)  
}
```

```
writerFinish () {  
    lock (RWLock)  
    numWriters--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

```
readerStart () {  
    lock (RWLock)  
    while (numWriters > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numReaders++  
    unlock (RWLock)  
}
```

```
readerFinish () {  
    lock (RWLock)  
    broadcast (rowfCV)  
    numReaders--  
    unlock (RWLock)  
}
```



Sure, both are protected by RWVLock.

Reader-writer code

```
writerStart () {  
    lock (RWLock)  
    while (numWriters > 0 ||  
           numReaders > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numWriters++;  
    unlock (RWLock)  
}
```

```
writerFinish () {  
    lock (RWLock)  
    numWriters--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

```
readerStart () {  
    lock (RWLock)  
    while (numWriters > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numReaders++  
    unlock (RWLock)  
}
```

```
readerFinish () {  
    lock (RWLock)  
    numReaders--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

If writer leaves with waiting readers and writers, who wins?

Reader-writer code

```
writerStart () {  
    lock (RWLock)  
    while (numWriters > 0 ||  
           numReaders > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numWriters++;  
    unlock (RWLock)  
}
```

```
writerFinish () {  
    lock (RWLock)  
    numWriters--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

```
readerStart () {  
    lock (RWLock)  
    while (numWriters > 0) {  
        wait (RWlock, rowfCV)  
    }  
    numReaders++  
    unlock (RWLock)  
}
```

```
readerFinish () {  
    lock (RWLock)  
    numReaders--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

How long might a writer have to wait?
(Don't want to starve the writers!)

Priority

- We want to give waiting writer threads a higher priority than waiting reader threads
- This will also help prevent writer starvation
- How can this be accomplished?
 - More complexity
 - Keep track of waiting and active writers separately
 - Only allow new reader access if there are no waiting or active writers
 - This might starve readers, but we are assuming that there are far more readers than writers!

Prioritizing waiting writers

```
writerStart () {  
    lock (RWLock)  
    while (actWriters > 0 ||  
           numReaders > 0) {  
        waitWriters++  
        wait (RWlock, rowfCV)  
        waitWriters--  
    }  
    actWriters++;  
    unlock (RWLock)  
}
```

```
writerFinish () {  
    lock (RWLock)  
    actWriters--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```

```
readerStart () {  
    lock (RWLock)  
    while ((actWriters +  
           waitWriters) > 0){  
        wait (RWlock, rowfCV)  
    }  
    numReaders++  
    unlock (RWLock)  
}
```

```
readerFinish () {  
    lock (RWLock)  
    numReaders--  
    broadcast (rowfCV)  
    unlock (RWLock)  
}
```