

ECE 570/670

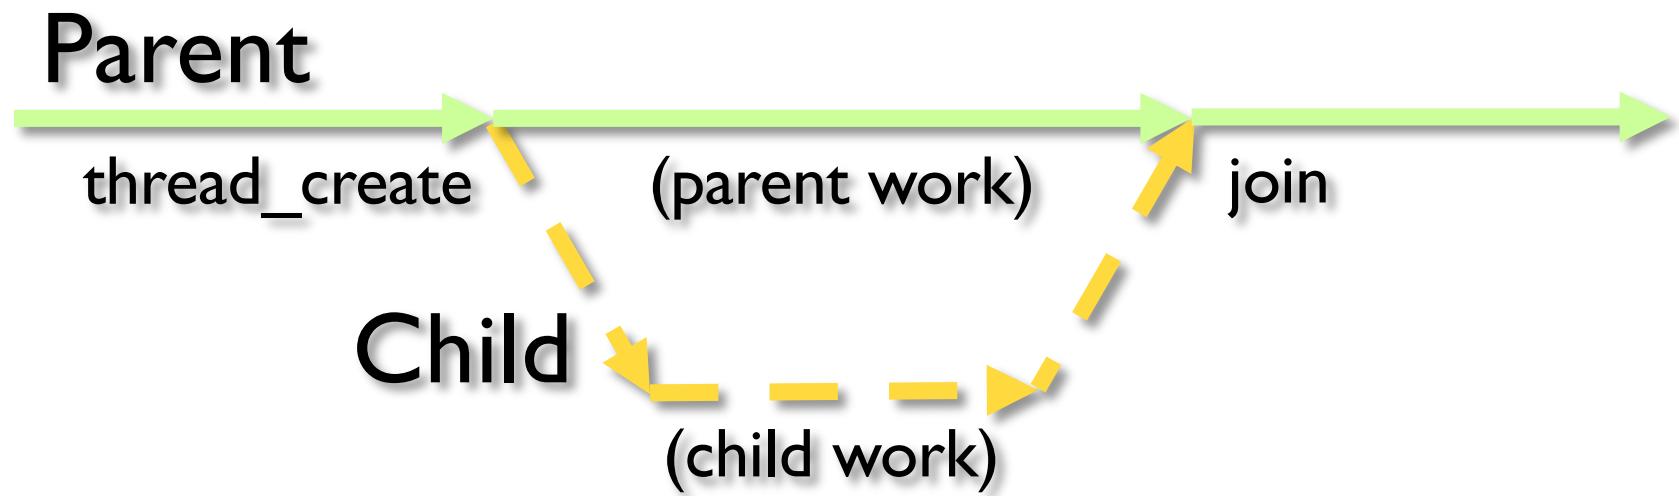
David Irwin
Lecture 9

Administrative Details

- 670 project meetings tomorrow
 - In my office: Knowles Eng. Bldg. 215C
- Should be working on Assignment 1d
 - Assignment 1 disk scheduler (1d)
 - How to use concurrency primitives
 - Assignment 1 thread library (1t)
 - How concurrency primitives actually work
 - Due 3/24
 - **Midterm in two weeks (3/29)**
 - Review in class on 3/28

Thread join

- How can the parent wait for child to finish?



Thread join

- Will this work?

- Sometimes, assuming
 - Uni-processor
 - No pre-emptions

- Never, ever assume these things!
- Yield is like slowing the CPU
 - Program must work +- any yields

```
child () {  
    print "child works"  
}  
  
parent () {  
    create child thread  
    print "parent works"  
    yield ()  
    print "parent continues"  
}
```

parent works
child works
parent continues
child works
parent works
parent continues

Thread join

- Will this work?

```
1 parent () {  
    create child thread  
    lock  
3    print "parent works"  
    wait  
    print "parent continues"  
    unlock  
}  
  
2 child () {  
    lock  
    print "child works"  
    signal  
    unlock  
}
```

- No. Child can call signal first.
- Would this work with semaphores?
 - Yes
 - No missed signals (increment sem value)

parent works
child works
parent continues
child works
parent works
parent continues

How can we solve this?

```
parent () {  
    childDone = 0  
    create child thread  
    print "parent works"  
    lock  
    while (!childDone) {  
        wait  
    }  
    print "parent continues"  
    unlock  
}
```

```
child () {  
    print "child works"  
    lock  
    childDone = 1  
    signal  
    unlock  
}
```

parent works
child works
parent continues

child works
parent works
parent continues

Thread join

- Many libraries supply “join” operation
 - Makes this a lot easier

```
parent () {  
    create child thread  
    print "parent works"  
    join with child  
    print "parent continues"  
}
```

```
child () {  
    print "child works"  
}
```

Switching threads

- What needs to happen to switch threads?

1. Thread returns control to OS

- For example, via the “yield” call

2. OS chooses next thread to run

3. OS saves state of current thread

- To its thread control block

4. OS loads context of next thread

- From its thread control block

5. Run the next thread

Assignment I
Part 2

I. Thread returns control to OS

- How does the thread return control back to OS?
 - Voluntary **internal events**
 - Thread might block inside lock or wait
 - Thread might call into kernel for service
 - (system call)
 - Thread might call yield
- Are internal events enough?

I. Thread returns control to OS

- Involuntary **external events**
 - (events not initiated by the thread)
 - Hardware interrupts
 - Transfer control directly to OS interrupt handlers
 - CPU checks for interrupts while executing
 - Jumps to OS code with interrupt mask set
 - Interrupts lead to pre-emption (a forced yield)
 - Common interrupt: timer interrupt

2. Choosing the next thread

- If no ready threads, just spin
 - Modern CPUs: execute a “halt” instruction
- Loop switches to thread if one is ready
- Many ways to prioritize ready threads
 - Involves “scheduling” the CPU
 - Will discuss next week

3. Saving state of current thread

- What needs to be saved?
 - Registers, PC, SP
- What makes this tricky?
 - Self-referential sequence of actions
 - Need registers to save state
 - But you're trying to save all the registers
- Saving the PC is particularly tricky

Saving the PC

- Why won't this work?

Instruction

address



100 store PC in TCB

101 switch to next thread

- Returning thread will execute instruction at 100
 - And just re-execute the switch
 - Really want to save address 102

4. OS loads the next thread

- Where is the thread's state/context?
 - Thread control block (in memory)
- How to load the registers?
 - Use load instructions to grab from memory
- How to load the stack?
 - Stack is already in memory, load SP

5. OS runs the next thread

- How to resume thread's execution?
 - Jump to the saved PC
- On whose stack are these steps running? (or who jumps to the saved PC?)
 - The thread that called yield
 - (or was interrupted or called lock/wait)
- How does this thread run again?
 - Some other thread must switch to it

Example thread switching

```
Thread 1
1  print "start thread 1"
2  yield ()
3  print "end thread 1"

Thread 2
4  print "start thread 2"
5  yield ()
6  print end thread 2"

yield ()
7  print "start yield (thread %d)"
8  swapcontext (tcb1, tcb2)
9  print "end yield (thread %d)"

swapcontext (tcb1, tcb2)
10 save regs to tcb1
11 load regs from tcb2
  // sp points to tcb2's stack now!
12 jump tcb2.pc
  // sp must point to tcb1's stack!
13 return
```

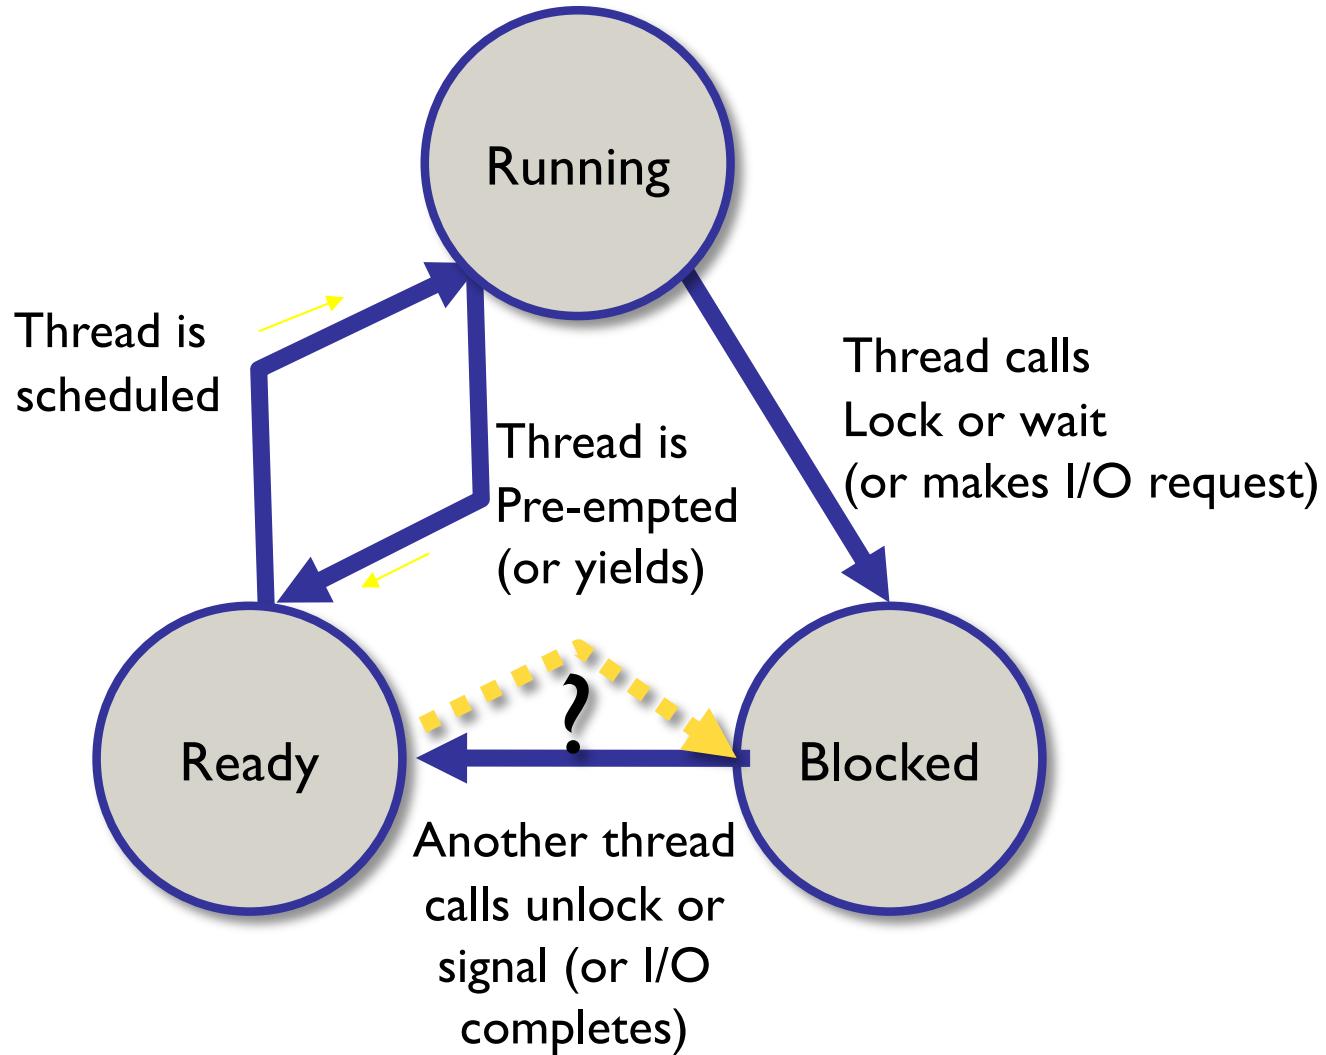
Thread 1 output	Thread 2 output
-----	-----
start thread 1	
start yield (thread 1)	
	start thread 2
	start yield (thread 2)
end yield (thread 1)	
end thread 1	
	end yield (thread 2)
	end thread 2

Note: this assumes no pre-emptions.
Programmers must assume pre-emptions, so
this output is not guaranteed.

Thread states

- Running
 - Currently using the CPU
- Ready
 - Ready to run, but waiting for the CPU
- Blocked
 - Stuck in lock (), wait () or down ()

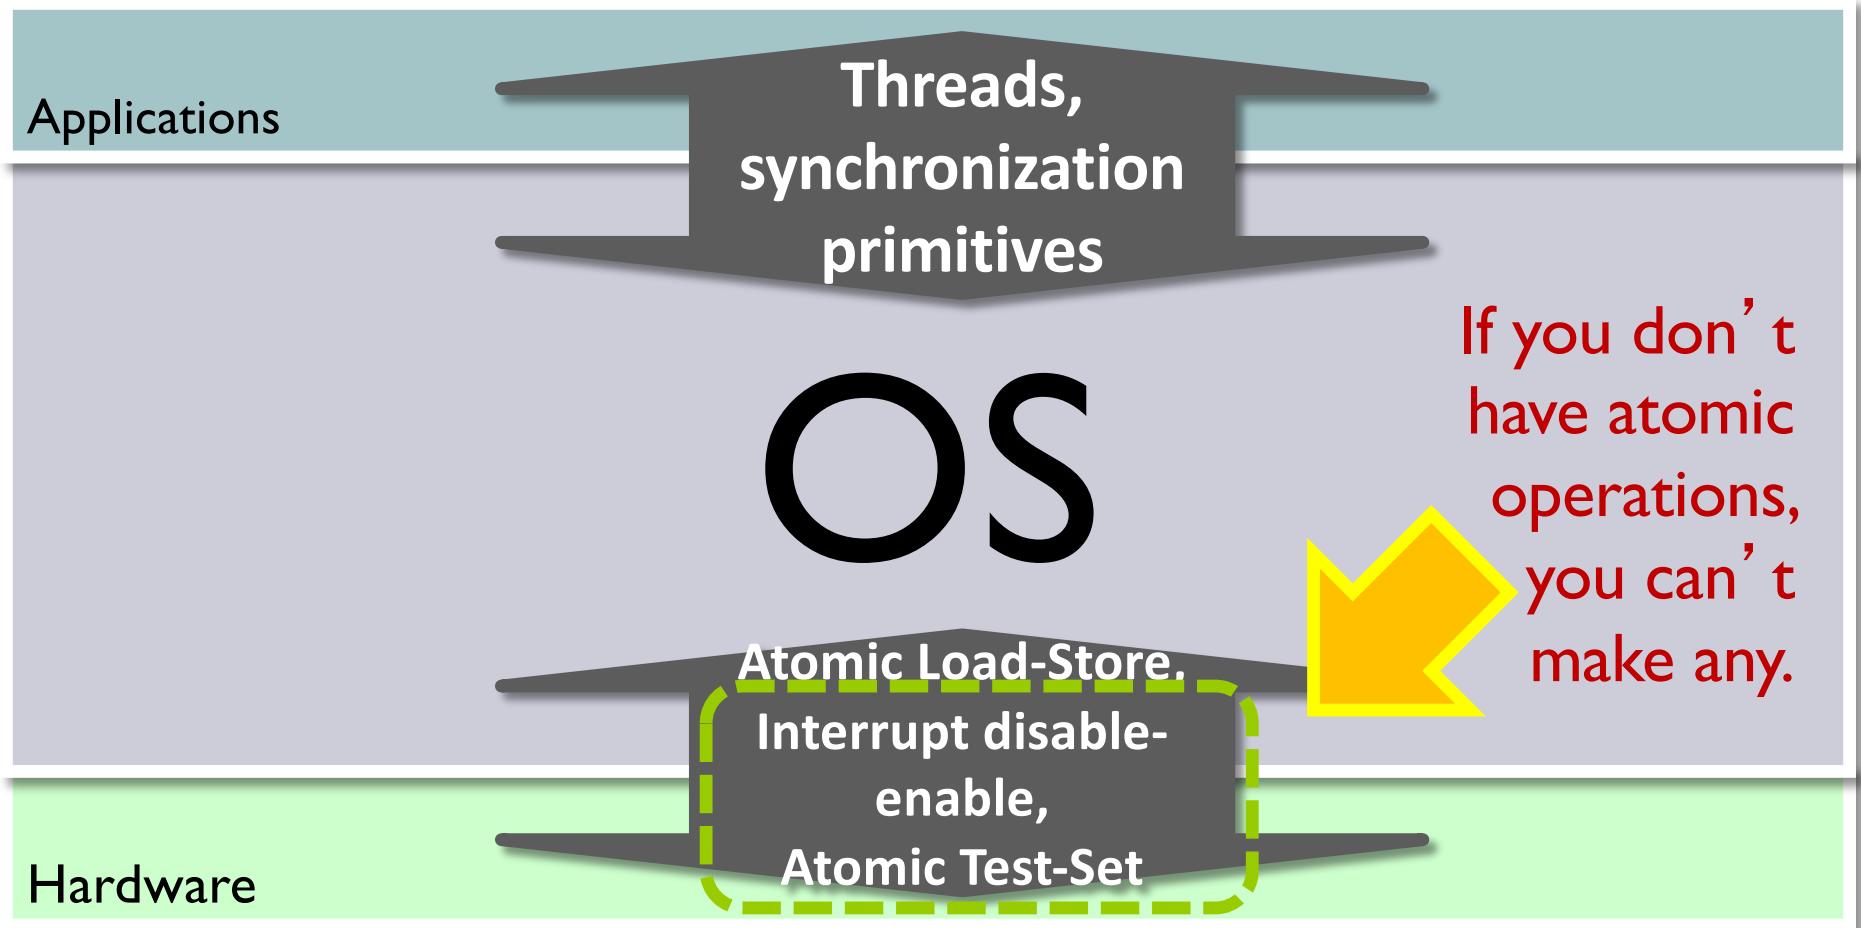
Thread states



Moving on...

- Discuss how locks are implemented

Recap and looking ahead



Using interrupt disable-enable

- Disable-enable on a uni-processor
 - Assume atomic (can use atomic load/store)
- How do threads get switched out (2 ways)?
 - Internal events (yield, I/O request)
 - External events (interrupts, e.g. timers)
- Easy to prevent internal events
- Use disable-enable to prevent external events

Why use locks?

- If we have disable-enable, why do we need locks?
 - Program could bracket critical sections with disable-enable
 - Might not be able to give control back to thread library

```
|-----|
| disable interrupts      |
| while (1){}           |
|-----|
```

- Can't have multiple locks (over-constrains concurrency)
- Assignment 1: only disable interrupts in thread library

Lock implementation #1

- Disable interrupts, busy-waiting

```
lock () {  
    disable interrupts  
    while (value != FREE) {  
        enable interrupts  
        disable interrupts  
    }  
    value = BUSY  
    enable interrupts  
}
```

```
unlock () {  
    disable interrupts  
    value = FREE  
    enable interrupts  
}
```

Why is it ok for lock code to disable interrupts?

It's in the trusted kernel (we have to trust something).

Lock implementation #1

- Disable interrupts, busy-waiting

```
lock () {  
    disable interrupts  
    while (value != FREE) {  
        enable interrupts  
        disable interrupts  
    }  
    value = BUSY  
    enable interrupts  
}
```

```
unlock () {  
    disable interrupts  
    value = FREE  
    enable interrupts  
}
```

Do we need to disable interrupts in unlock?

Only if “value = FREE” is multiple instructions (safer)

Lock implementation #1

- Disable interrupts, busy-waiting

```
lock () {  
    disable interrupts  
    while (value != FREE) {  
        enable interrupts  
        disable interrupts  
    }  
    value = BUSY  
    enable interrupts  
}
```

```
unlock () {  
    disable interrupts  
    value = FREE  
    enable interrupts  
}
```

Why enable-disable in lock loop body?

Otherwise, no one else will run (including unlockers)

Using read-modify-write instructions

- Disabling interrupts
 - Ok for uni-processor, breaks on multi-processor
- Could use atomic load-store to make a lock
 - Inefficient, lots of busy-waiting
- *Need hardware support!*

Using read-modify-write instructions

- Most modern processor architectures
 - Provide an atomic read-modify-write instruction
- Atomically
 - Read value from memory into register
 - Write new value to memory
- Implementation details
 - Lock memory location at the memory controller

Too much milk, Solution 2

```
if (noMilk) {  
    if (noNote){  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

Block is not atomic.
Must atomically

- check if lock is free
- grab it

What key part of lock code had to be atomic?

Test&set on most architectures

```
test&set (X) {  
    tmp = X  
    X = 1  
    return (tmp)  
}
```

Set: sets location to 1
Test: returns old value



- Slightly different on Intel x86 (Exchange)
 - Atomically swaps value between register and memory

Lock implementation #2

- Use test&set
- Initially, value = 0

```
lock () {  
    while (test&set(value) == 1) {  
    }  
}
```

```
unlock () {  
    value = 0  
}
```

What happens if value = 1?
What happens if value = 0?

Locks and busy-waiting

- All previous implementations have used busy-waiting
 - Wastes CPU cycles
 - Better for thread to sleep and let other threads run
- To reduce busy-waiting, integrate
 - Lock implementation
 - Thread dispatcher data structures

Lock implementation #3

- Interrupt disable, no busy-waiting

```
lock () {
    disable interrupts
    if (value == FREE) {
        value = BUSY // lock acquire
    } else {
        add thread to queue of threads waiting for lock
        switch to next ready thread // don't add to ready queue
    }
    enable interrupts
}

unlock () {
    disable interrupts
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    enable interrupts
}
```

Lock implementation #3

This is
called a
“hand-off”
lock.

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        add thread to queue of threads waiting for lock  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```

Who gets the lock after someone calls unlock?

Thread calling unlock gives lock to thread on wait queue

Lock implementation #3

This is
called a
“hand-off”
lock.

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        add thread to queue of threads waiting for lock  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```

Who might get the lock if it weren't handed-off directly? (e.g. if value weren't set BUSY in unlock)

Lock implementation #3

This is
called a
“hand-off”
lock.

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        add thread to queue of threads waiting for lock  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```

What kind of ordering of lock acquisition guarantees does the hand-off lock provide?

Lock implementation #3

This is
called a
“hand-off”
lock.

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        add thread to queue of threads waiting for lock  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```



What does this mean?

Lock implementation #3

This is
called a
“hand-off”
lock.

```
lock () {
    disable interrupts
    if (value == FREE) {
        value = BUSY // lock acquire
    } else {
        lockqueue.push(&current_thread->ucontext);
        swapcontext(&current_thread->ucontext,
                    &new_thread->ucontext));
    }
    enable interrupts
}

unlock () {
    disable interrupts
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    enable interrupts
}
```



Adding a pointer to the TCB/context.

Lock implementation #3

This is
called a
“hand-off”
lock.

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        add thread to queue of threads waiting for lock  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```

Why a separate queue for the lock?

All locked threads are not ready to run; can't control ordering

Lock implementation #3

This is
called a
“hand-off”
lock.

```
lock () {
    disable interrupts
    if (value == FREE) {
        value = BUSY // lock acquire
    } else {
        add thread to queue of threads waiting for lock
        switch to next ready thread // don't add to ready queue
    }
    enable interrupts
}

unlock () {
    disable interrupts
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    enable interrupts
}
```

Assignment 1 note: you must guarantee FIFO ordering
of lock acquisition.

Lock implementation #3

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        enable interrupts  
        add thread to queue of threads waiting for lock  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```



When and how could this fail?

Someone can unlock before thread added to wait queue

Lock implementation #3

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        add thread to queue of threads waiting for lock  
        enable interrupts  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```

1

3

2

When and how could this fail?

Lock implementation #3

```
lock () {
    disable interrupts
    if (value == FREE) {
        value = BUSY // lock acquire
    } else {
        add thread to queue of threads waiting for lock
        switch to next ready thread // don't add to ready queue
    }
    enable interrupts
}

unlock () {
    disable interrupts
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    enable interrupts
}
```

Putting lock thread on lock wait queue, switch must be atomic. Must call switch with interrupts off.