

# ECE 570/670

David Irwin  
Lecture 10

# Administrative Details

- Should be working on Assignment I
  - Assignment I disk scheduler (Id)
    - How to use concurrency primitives
  - Assignment I thread library (It)
    - How concurrency primitives actually work
  - **Midterm coming up (3/29)**
    - Review in class on 3/28
- Piazza questions are great
- Will post videos in advance of mid-term

# Administrative Details

- ECE670 Project Proposals (due Friday, 3/24)
  - Spoke with most groups
  - Other groups should drop by office hours

# Lock implementation #3

```
lock () {  
    disable interrupts  
    if (value == FREE) {  
        value = BUSY // lock acquire  
    } else {  
        add thread to queue of threads waiting for lock  
        switch to next ready thread // don't add to ready queue  
    }  
    enable interrupts  
}  
  
unlock () {  
    disable interrupts  
    value = FREE  
    if anyone on queue of threads waiting for lock {  
        take waiting thread off queue, put on ready queue  
        value = BUSY  
    }  
    enable interrupts  
}
```

Putting lock thread on lock wait queue, switch must be atomic. Must call switch with interrupts off.

# How is switch returned to?

- Recall swapcontext
- Think of switch as three phases
  - Save current location (SP, PC)
  - Point processor to another stack ( $SP'$ )
  - Jump to another instruction ( $PC'$ )
- Only way to get back to a switch?
  - Have **another** thread call switch

# Lock implementation #3

```
lock () {
    disable interrupts
    if (value == FREE) {
        value = BUSY // lock acquire
    } else {
        add thread to queue of threads waiting for lock
        switch to next ready thread // don't add to ready queue
    }
    enable interrupts
}

unlock () {
    disable interrupts
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    enable interrupts
}
```

What is lock() assuming about the state of interrupts after switch returns?

# Interrupts and returning to switch

- Lock() can assume that switch
  - Is always called with interrupts disabled
- On return from switch
  - Previous thread must have disabled interrupts
- Next thread to run
  - Becomes responsible for re-enabling interrupts
- **Invariants:** threads promise to
  - Disable interrupts before switch is called
  - Re-enable interrupts after returning from switch

## Thread A

```
enable interrupts  
} // exit thread library function  
<user code>  
lock () {  
    disable interrupts  
    ...  
    switch  
  
back from switch  
...  
enable interrupts  
} // exit lock  
<user code>
```

## Thread B

```
yield () {  
    disable interrupts  
    ...  
    switch  
  
    back from switch  
    ...  
    enable interrupts  
} // exit yield  
<user code>  
unlock () // moves A to ready queue  
yield () {  
    disable interrupts  
    ...  
    switch
```



# Lock implementation #4

- Test&set, minimal busy-waiting

```
lock () {
    while (test&set (guard)) {} // like interrupt disable
    if (value == FREE) {
        value = BUSY
    } else {
        put on queue of threads waiting for lock
        switch to another thread // don't add to ready queue
    }
    guard = 0 // like interrupt enable
}

unlock () {
    while (test&set (guard)) {} // like interrupt disable
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    guard = 0 // like interrupt enable
}
```

# Lock implementation #2

- Use test&set
- Initially, value = 0

```
lock () {  
    while (test&set(value) == 1) {  
    }  
}
```

```
unlock () {  
    value = 0  
}
```

What happens if value = 1?  
What happens if value = 0?

# Lock implementation #4

Why is this better than implementation #2?

Only busy-wait while another thread is in lock or unlock

Before, we busy-waited while lock was held

```
lock () {
    while (test&set (guard)) {} // like interrupt disable
    if (value == FREE) {
        value = BUSY
    } else {
        put on queue of threads waiting for lock
        switch to another thread // don't add to ready queue
    }
    guard = 0 // like interrupt enable
}

unlock () {
    while (test&set (guard)) {} // like interrupt disable
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    guard = 0 // like interrupt enable
}
```

# Lock implementation #4

What is the switch invariant?

Threads promise to call switch with guard set to 1.

```
lock () {
    while (test&set (guard)) {} // like interrupt disable
    if (value == FREE) {
        value = BUSY
    } else {
        put on queue of threads waiting for lock
        switch to another thread // don't add to ready queue
    }
    guard = 0 // like interrupt enable
}

unlock () {
    while (test&set (guard)) {} // like interrupt disable
    value = FREE
    if anyone on queue of threads waiting for lock {
        take waiting thread off queue, put on ready queue
        value = BUSY
    }
    guard = 0 // like interrupt enable
}
```

# Summary of implementing locks

- Synchronization code needs atomicity
- Three options
  - **Atomic load-store**
    - Lots of busy-waiting
  - **Interrupt disable-enable**
    - No busy-waiting
    - Breaks on a multi-processor machine
  - **Atomic test-set**
    - Minimal busy-waiting
    - Works on multi-processor machines

# Assignment I (Part 2)

- Write the thread library used in Part I
  - Use Linux calls provided for building user-level thread libraries
    - `getcontext()`; `makecontext()`; `swapcontext()`;
    - Operate on `ucontext` structs (i.e., the TCB)
    - Example of creating new threads on website
    - Also, see man pages
  - Given `libinterrupt` library
    - Allows you to enable and disable interrupts in your thread library
    - Only disable interrupts in your thread library code (not in applications!)
  - Look at pseudocode from class
    - Implement ready queue, wait queue for lock, and wait queue for signal
    - Locks should be acquired in order they are requested (FIFO)

# Assignment I (Part 2)

- Write the thread library used in Part I
  - `getcontext();`
    - Initialization of a context, e.g., TCB
  - `makecontext();`
    - Takes function associated with the context as argument
    - Should have this point to a function **startFunction** *in your thread library*
      - This function takes the user function as an argument
    - **startFunction** will call the user's thread function
      - When user function returns the thread is finished executing
  - `swapcontext();`
    - As in the slides

# Assignment I (Part 2)

- Error Handling
  - Must detect library misuse
    - E.g., calling thread function before calling `thread_libinit`, calling `thread_libinit` multiple times, etc.
  - Must detect if functions it calls return an error
    - Return -1 from thread function on errors
  - Use “asserts” throughout thread library to catch logic errors in your thread library
- Getting started
  - Start with implementing `thread_libinit`, `thread_create`, and `thread_yield`
  - After those working, then work on the monitor functions

# Assignment I (Part 2)

- You will also need to work on thread library test suite
- For each of your test cases
  - Compare output of your library with `thread.o`
- Writing test cases
  - Read through spec
    - Write test cases to stress each required part
      - E.g. `lock()` blocks if lock is already held
    - Use `yield()` to create the right inter-leavings
  - Read through your code
    - Write test cases to exercise all lines of code
    - E.g. each clause of an `if/else` statement

# Assignment I (Part 2)

- Micro-tests better for debugging than macro-tests
  - One test case usually uncovers many bugs
  - Use disk scheduler and example program as templates for creating test cases
- Test cases will be graded on how thoroughly they exercise your thread library
  - Will run them against thread libraries with various common bugs and see if you test case exposes those bugs

# Assignment I (Part 2)

- Compiling and running
  - Must be in a single file thread.cc
  - **g++ thread.cc app.cc libinterrupt.a -ldl -o app**
  - Thread library only generates one output line
    - “Thread library exiting.”
  - **submit570 lt thread.cc test1.cc test2.cc**

# Hints on Part 2

- Warnings are probably ok; errors are not ok
- Make sure you implement hand-off locks, as discussed in class
- Make sure you can create lots of threads without running out of memory
- Make sure you delete threads that are finished
  - How do you know a thread finishes?
- Make sure you enable and disable interrupts appropriately
  - Follow invariants discussed in class
- Test cases: not just for the autograder, mostly for your own code!

# Concurrency so far

- Mostly tried to constrain orderings
  - Locks, monitors, semaphores
- Note that it is possible to over-constrain too
  - A must run before B
  - B must run before A
- This can lead to **deadlock**

# Definitions

- **Resource**
  - “Thing” needed by a thread to do its job
  - Threads wait for resources
  - “Things” can be locks, disk space, memory, CPU
- **Deadlock**
  - Circular waiting for resources
  - Leaves all waiting threads unable to make progress

# Example deadlock

- Both CS187 and ECE241 are full
  - You want to switch from 187 to 241
  - Someone else wants to switch from 187 to 241
- Algorithm for switching
  - Wait for spot in new class to open up
  - Add new class
  - Drop old class
- Problem: must add before dropping

# Deadlock and starvation

- Deadlock → starvation
  - Starvation = one process waits forever

## Thread A

```
lock (x)  
lock (y)  
...  
unlock (y)  
unlock (x)
```

## Thread B

```
lock (y)  
lock (x)  
...  
unlock(x)  
unlock(y)
```

Can deadlock occur?

# Deadlock and starvation

- Deadlock → starvation
  - Starvation = one process waits forever

**Thread A**

```
lock (x)  
  
lock (y) // wait for B
```

**Thread B**

```
lock (y)  
  
lock (x) // wait for A
```

Will deadlock occur?

# Common thread work pattern

```
Phase 1. while (not done) {  
    get some resources (block if necessary)  
    work // assume finite  
}  
Phase 2. release all resources
```

- This is called “two-phase” locking
  - Common in databases, but...
  - ...deadlock-prone

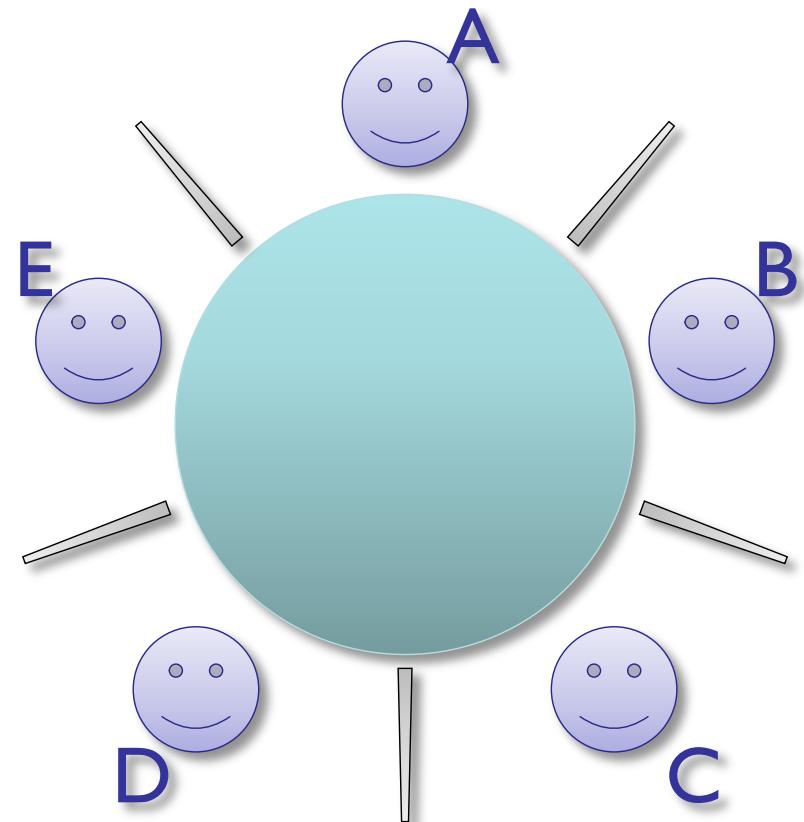
# Dining philosophers

Philosophers alternate between eating (noodles) and thinking

## Philosopher algorithm

- 1) Wait for right chopstick
- 2) Pick up right chopstick
- 3) Wait for left chopstick
- 4) Pick up left chopstick
- 5) Eat
- 6) Put both chopsticks down

How can deadlock occur?



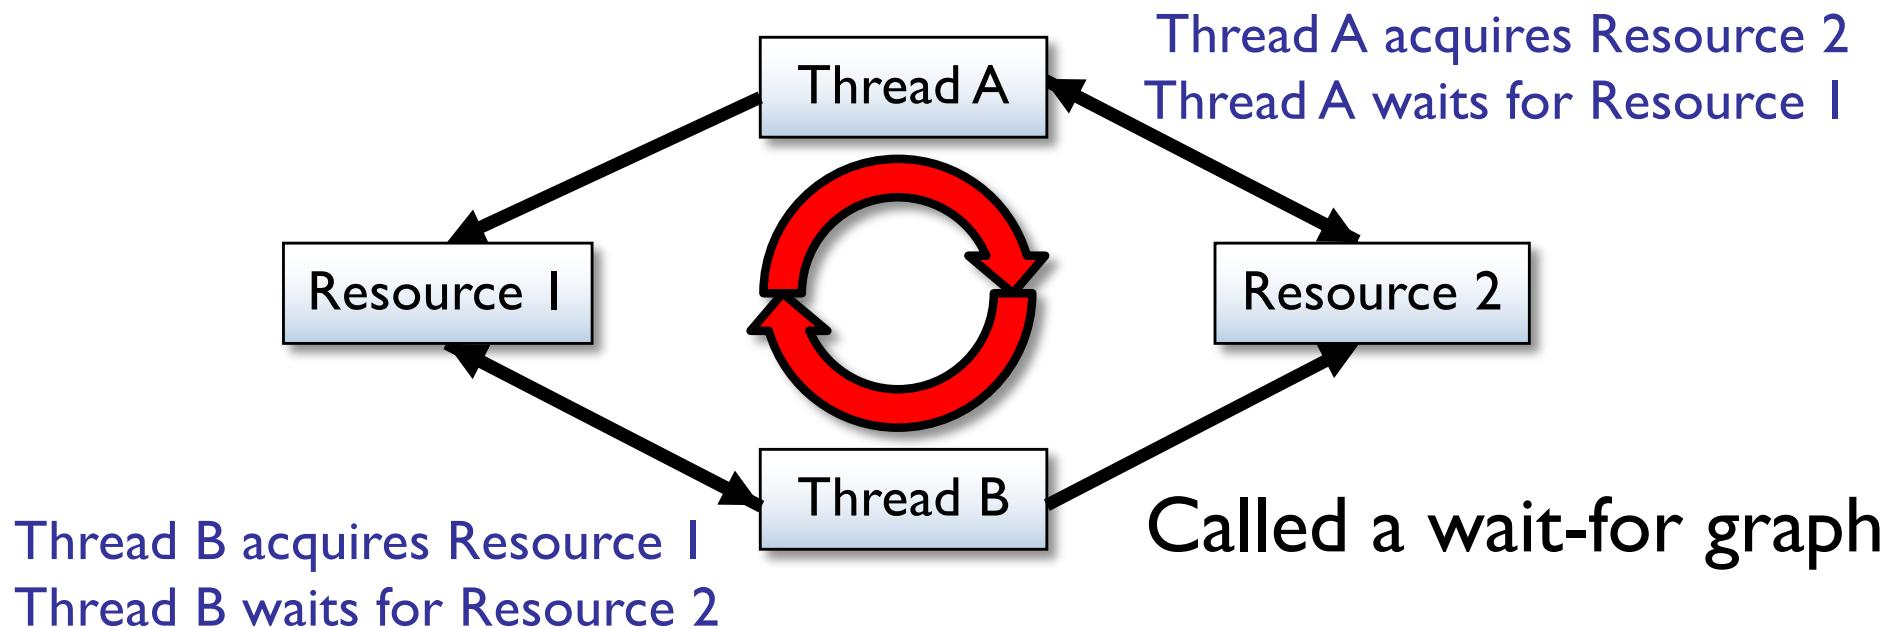
# Conditions for deadlock

1. Limited resource
  - Not enough for all threads simultaneously
2. Hold-and-wait
  - Hold one resource, while waiting for another
3. No pre-emption
  - Cannot force threads to give up resources
4. Circular chain of requests

# Circular chain of requests

- Arrows

- Thread → resource it's waiting for
- Resource → thread that's holding it



# Deadlock

- As the OS, what should we do about deadlocks?

## 1. Ignore them

- Common OS-level solution for programs
- (OS has no control over user programs)

## 2. Detect-and-fix

## 3. Prevent

# Detect-and-fix

- First part: **detect**
  - This is easy; just scan the wait-for graph
- Second part: **fix**
  1. Kill first thread, take back resources by force
    - Why might this be unsafe?
      - Can expose inconsistent state – invariants might be broken
  2. Roll-back actions of 1 or more thread, retry
    - Often used in databases
    - Not always possible (some actions can't be undone)
      - E.g. can't unsend a network message

# Detect-and-fix

- Retrying during fix phase can be tricky
  - If holding R and waiting for L, drop R and try again.

What could happen?

Everyone picks up R

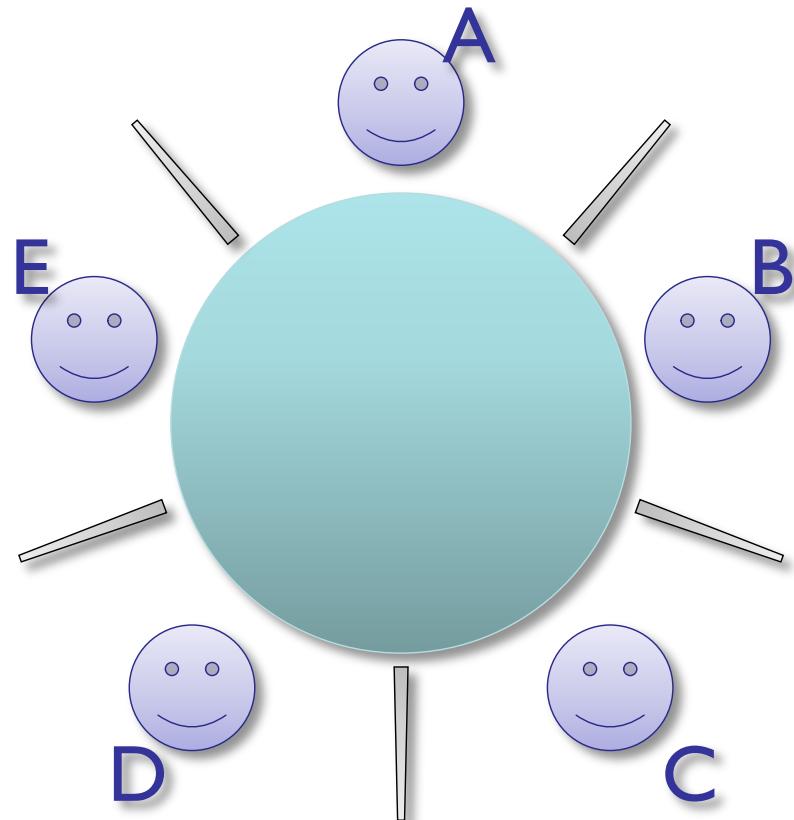
Everyone drops R

Everyone picks up R

Everyone drops R

(and so on)

This is called “livelock.”



# Detect-and-fix

- Retrying during fix phase can be tricky
  - If holding R and waiting for L, drop R and try again.

How to prevent livelock?

Choose a winner.

How to choose a winner?

First to start/pick up?

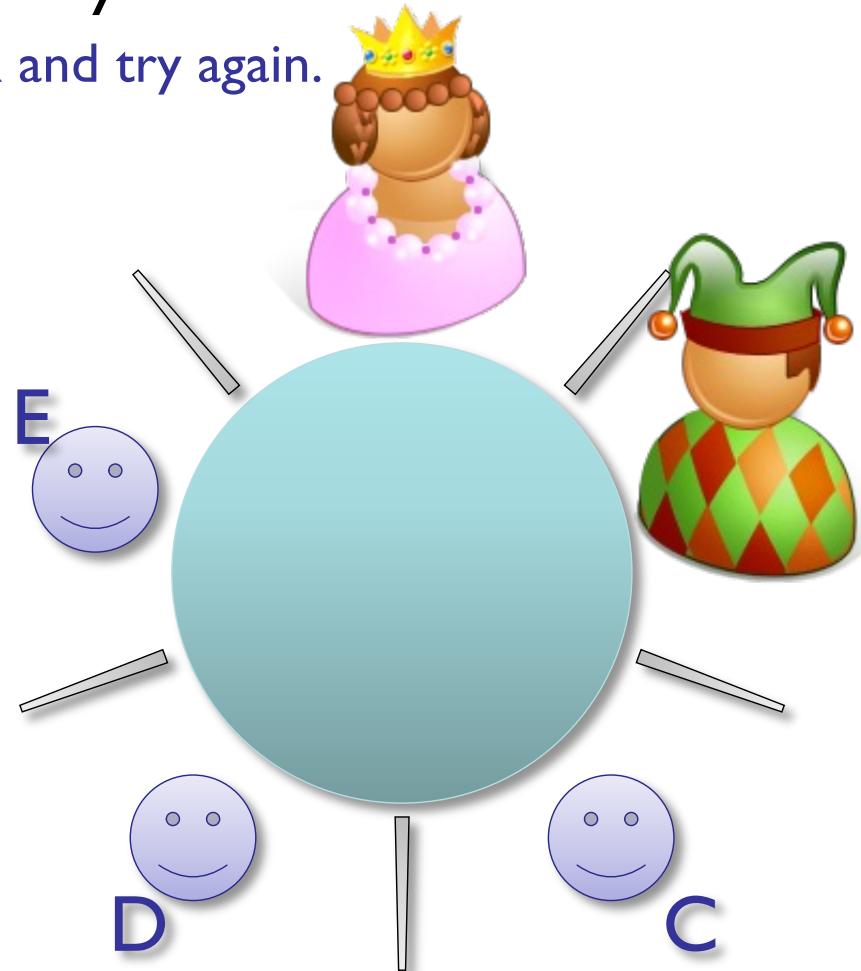
What if we have priorities?

(called “priority inversion”)

e.g. Queen waiting for joker

Should Queen always win?

Depends, could starve the joker



# Deadlock prevention

- Idea: remove one of four pre-conditions

## I. Make resources unlimited

- Best solution, but often impossible
- Can still increase number, reduce likelihood
- (larger 187 and 241 classes)
- This doesn't work for locks

# Deadlock prevention

## 2. Eliminate hold-and-wait

- Idea: move resource acquisition to beginning

```
Phase 1a. get all resources  
Phase 1b. while (not done) {  
            work // assume finite  
        }  
Phase 2. release all resources
```

- Two ways to avoid holding while waiting

# Eliminating hold-and-wait

- I. Wait for everything to be free, then grab them all at once

Philosopher algorithm

```
lock
while (left chopstick busy ||
       right chopstick busy) {
    wait
}
pick up right chopstick
pick up left chopstick
unlock
eat
lock
drop chopsticks
unlock
```

Any potential problems?

Could induce starvation (neighbors alternate eating, one starves).

# Eliminating hold-and-wait

2. If you find a resource busy, drop everything, and try again

```
Philosopher algorithm
```

```
lock
while (1) {
    while (right chopstick busy) {
        wait
    }
    pick up right chopstick
    if (left chopstick busy) {
        drop right chopstick
    } else {
        pick up left chopstick
        break
    }
}
unlock
eat
lock
drop chopsticks
unlock
```

Issues?

Overconstraints concurrency  
Hold reservations longer than needed.

# Deadlock prevention

## 3. Enable pre-emption (force threads to give up resources)

- Can pre-empt the CPU (context switch)
- Can pre-empt memory (swap to disk)
- Not everything can be pre-empted
- Can locks be pre-empted?
  - Probably not
  - Must ensure that data is in consistent state

# Deadlock prevention

## 4. Eliminate circular chain of requests

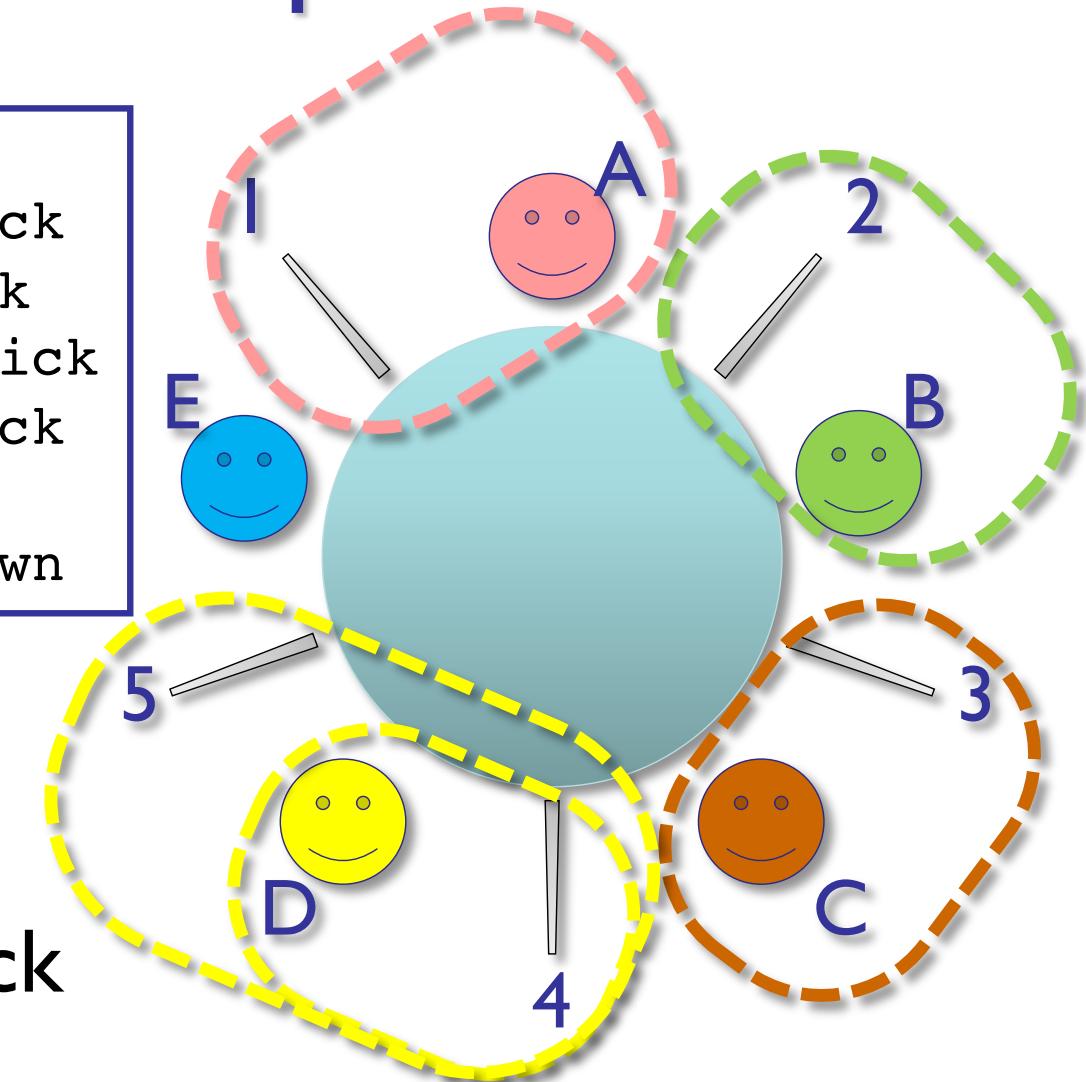
- How can we get rid of these cycles?
  - Impose an ordering on resource acquisition
  - E.g. must always acquire A before B
- In dining philosophers?
  - Number chopsticks
  - Pick up lower-numbered chopstick first

# Dining philosophers

## Philosopher algorithm

- 1) Wait for lower chopstick
- 2) Pick up lower chopstick
- 3) Wait for higher chopstick
- 4) Pick up higher chopstick
- 5) Eat
- 6) Put both chopsticks down

Try to create a deadlock



# Universal ordering

- Why does this work?
- At some point in time
  - T holds highest-numbered acquired resource
  - T is guaranteed to make progress. Why?
    - If T needs a higher-numbered resource, it must be free
    - If T needs a lower-numbered resource, it already has it
  - If T can make progress, it will eventually release all of its resources
  - What if another thread acquires a higher-numbered resource?
    - They just become T
    - (in which case, the same reasoning as above holds)