

# ECE 570/670

David Irwin  
Lecture 11

# Administrative Details

- Should be working on Assignment I
  - Assignment I disk scheduler (Id)
    - How to use concurrency primitives
    - Due today
  - Assignment I thread library (It)
    - How concurrency primitives actually work
  - **Midterm next Wednesday**
    - **Review next class**
- Will post example questions on calendar
  - Will go over answers in class
- Posted videos on Moodle for review (see Echo360 link)

# Administrative Details

- ECE670 project proposal due tomorrow
  - Upload to Moodle
  - A few groups have still not met with me

# Switching threads

- What needs to happen to switch threads?

## I. Thread returns control to OS

- For example, via the “yield” call

2. **OS chooses next thread to run**

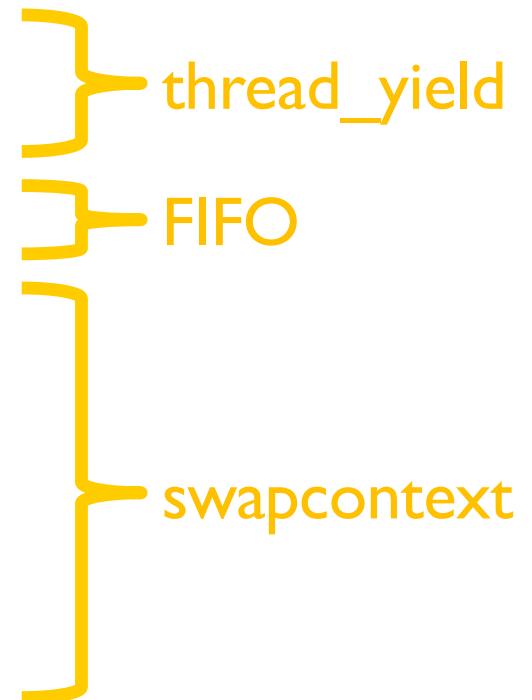
3. **OS saves state of current thread**

- To its thread control block

4. **OS loads context of next thread**

- From its thread control block

5. **Run the next thread**



# Scheduling goals

## I. **Minimize average response time**

- Elapsed time to do a job (what users care about)
- Try to maximize idle time
  - Incoming jobs can then finish as fast as possible

## 2. **Maximize throughput**

- Jobs per second (what admins care about)
- Try to keep parts as busy as possible
- Minimize wasted overhead (e.g. context switches)

# Scheduling goals

## 3. Fairness

- Share CPU among threads equitably
- Key question: what does “fair” mean?

Job 1

Needs 100  
seconds of  
CPU time

Job 2

Needs 100  
seconds of  
CPU time

# What does “fair” mean?

- How can we schedule these jobs?
    - Job 1 then Job 2
      - 1's response time = 100, 2's response time = 200
      - Average response time = 150
    - Alternating between 1 and 2
      - 1's response time = 2's response time = 200
      - Average response time = 200
- |   |   |
|---|---|
| <b>Job 1</b><br>Needs 100<br>seconds of<br>CPU time | <b>Job 2</b><br>Needs 100<br>seconds of<br>CPU time |
|---|---|

# Fairness

- First time thinking of OS as government
- Fairness can come at a cost
  - (in terms of average response time)
- Finding a balance can be hard
  - What if you have 1 big job, 1 small job?
  - Response time proportional to size?
  - Trade-offs come into play

# FCFS (first-come-first-served)

- FIFO ordering between jobs
- No pre-emption (run until done)
  - Run thread until it blocks or yields
  - No timer interrupts
- Essentially what you're doing in Project 1t

# FCFS example

- Job A (100 seconds), Job B (1 second)

A arrives and starts

( $t=0$ )



B arrives  
( $t=0+$ )

A done  
( $t=100$ )



B done  
( $t=101$ )



- Average response time = 100.5 seconds

# FCFS

- Pros
  - Simplicity
- Cons?
  - Short jobs stuck behind long ones
  - Non-interactive
  - (for long CPU job, user can't type input)

# Round-robin

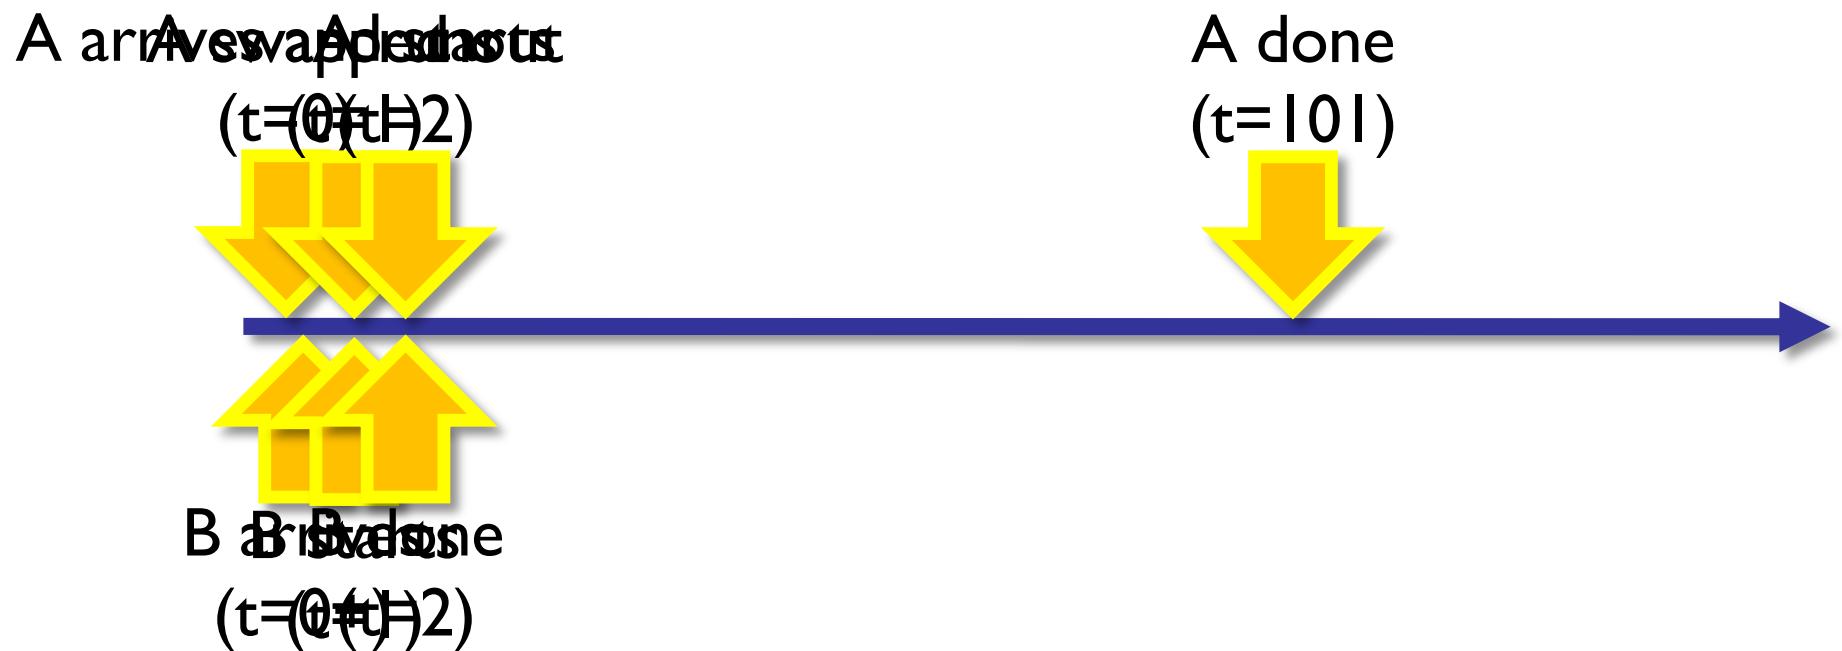
- Goal
  - Improve average response time for short jobs
- Solution to FCFS's problem
  - Add pre-emption!
  - Pre-empt CPU from long-running jobs
  - (timer interrupts)
- This is what most OSes do

# Round-robin

- In what way is round robin fairer than FCFS?
  - Makes response times ~ job length
- In what way is FCFS fairer than round robin?
  - First to start will have better response time
- Like two children with a toy
  - Should they take turns?
    - “She’s been playing with it for a long time.”
  - Should first get toy until she’s bored?
    - “I had it first.”

# Round-robin example

- Job A (100 seconds), Job B (1 second)

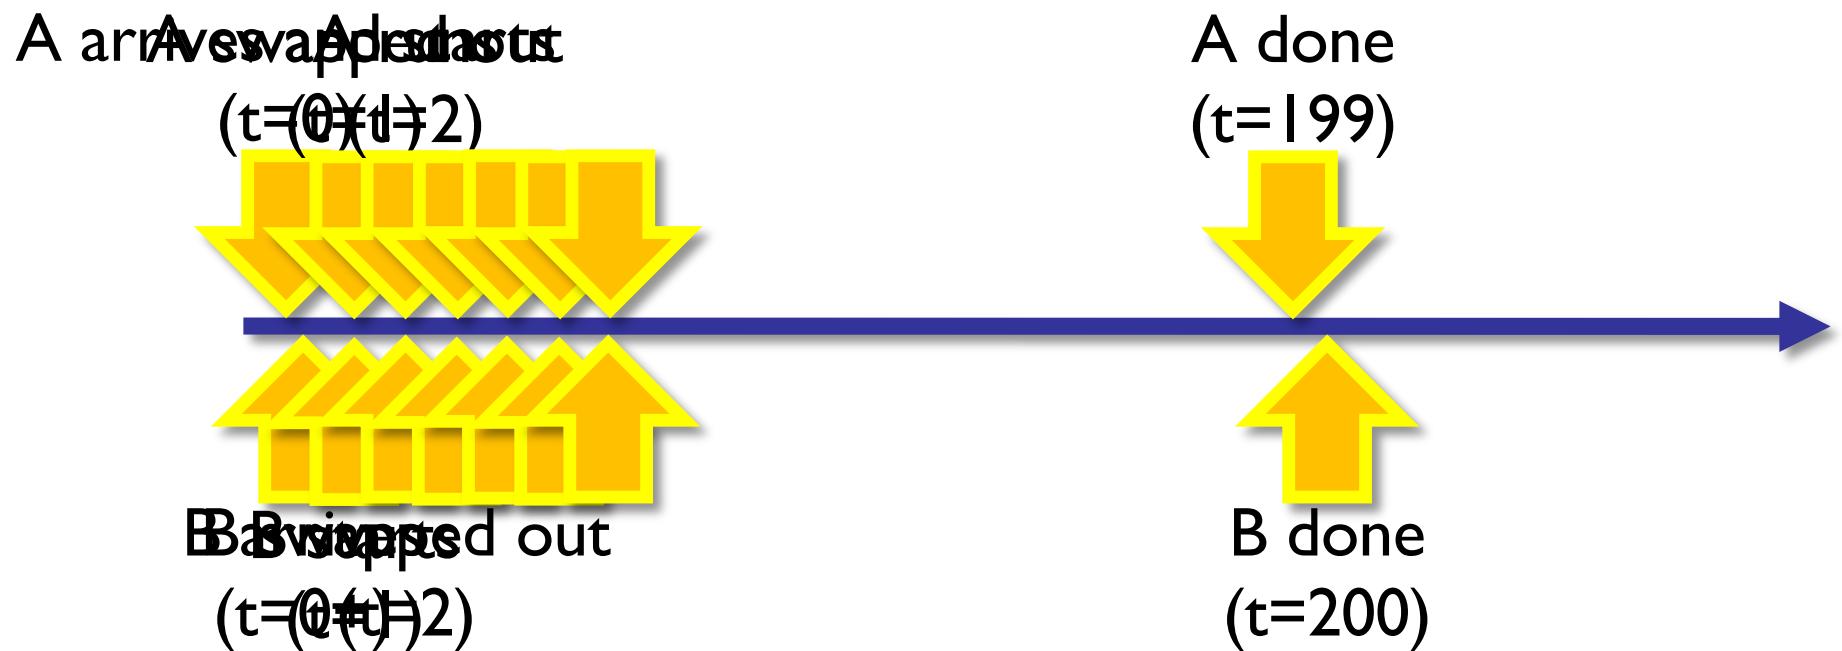


- Average response time = 51.5 seconds

Does round robin always provide lower response times than FCFS?

# Round-robin example 2

- Job A (100 seconds), Job B (100 second)

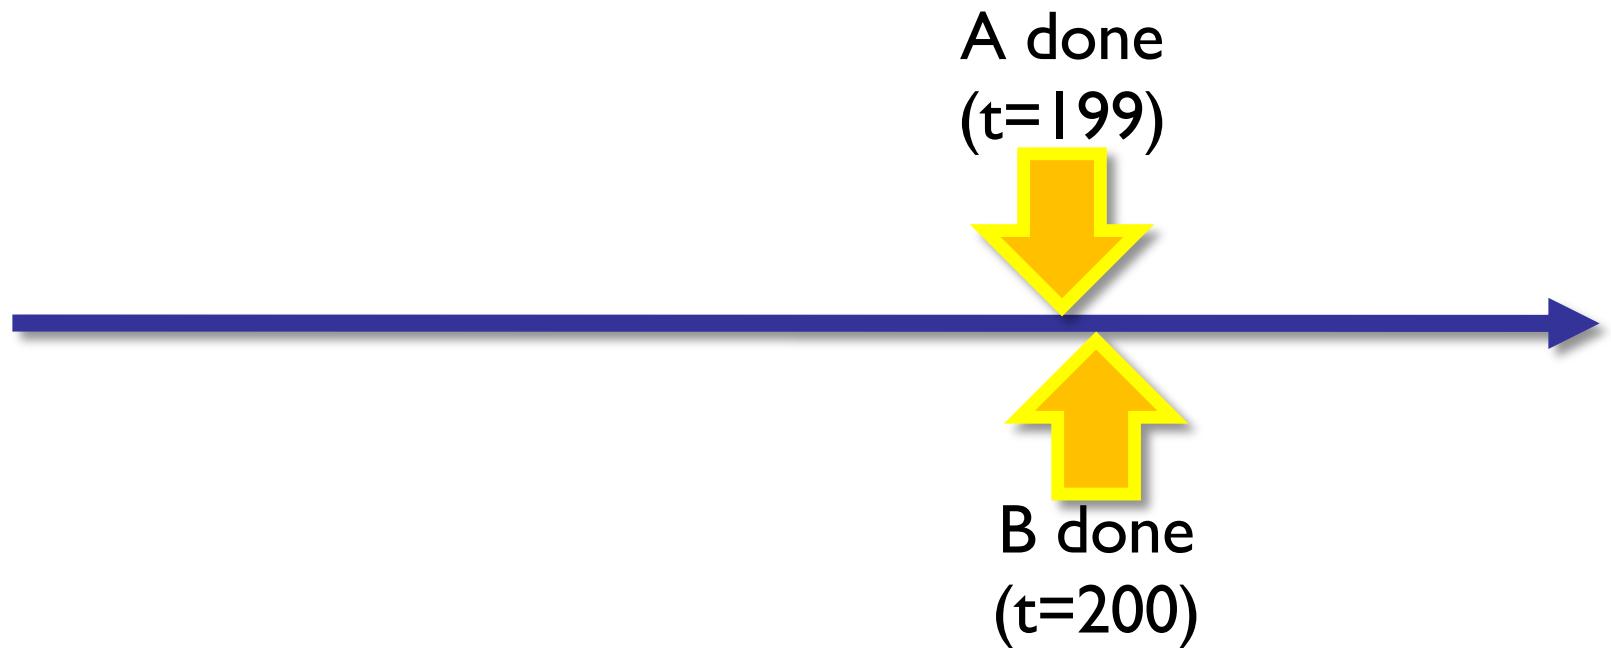


- Average response time = 199.5 seconds

Any hidden costs that we aren't counting? **Context switches**

# Round-robin example 2

- Job A (100 seconds), Job B (100 second)

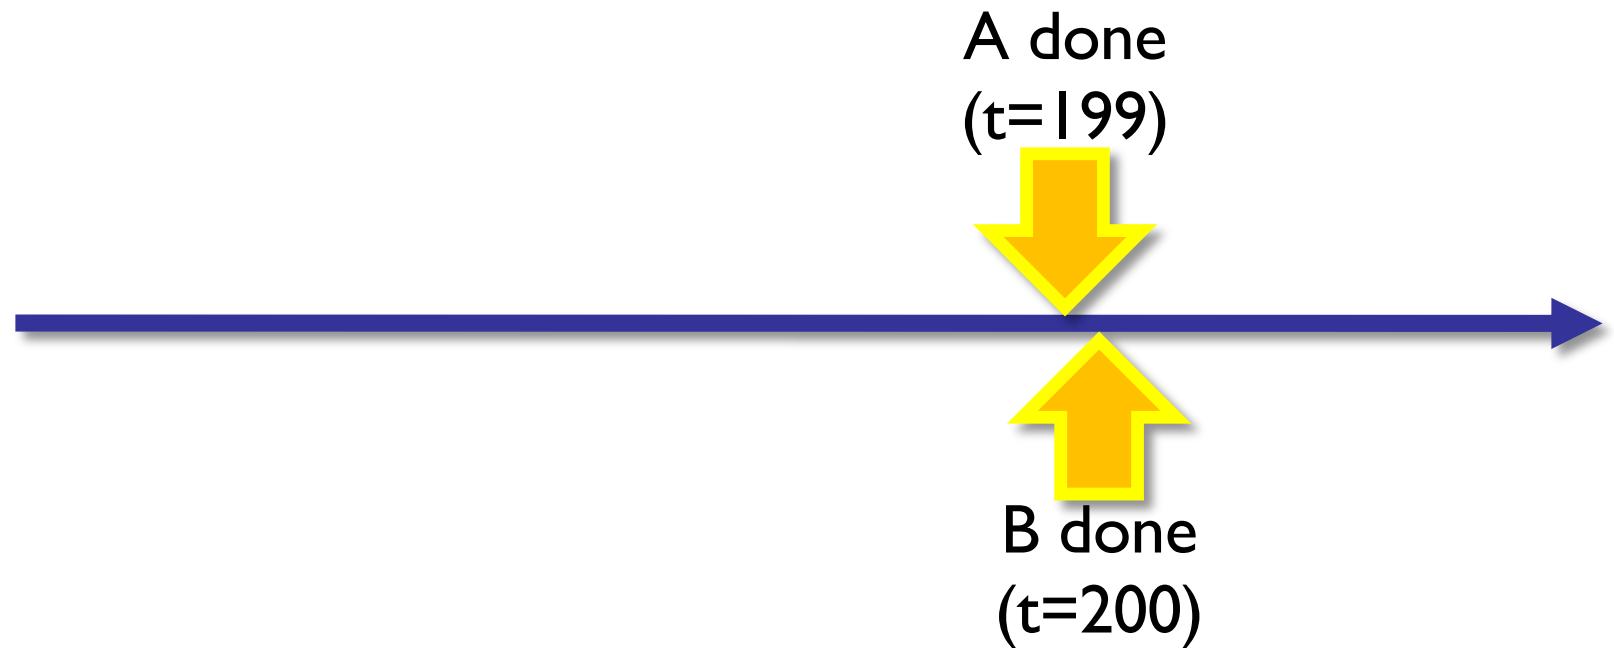


- Average response time = 199.5 seconds

What would FCFS's avg response time be? 150 seconds

# Round-robin example 2

- Job A (100 seconds), Job B (100 second)



- Average response time = 199.5 seconds

Which one is fairer? It depends ...

# Round-robin

- Pros
  - Good for interactive computing
  - Better than FCFS for mix of job lengths
- Cons
  - Less responsive for uniform job lengths
  - Hidden cost: context switch overhead
- How should we choose the time-slice?
  - Typically a compromise, e.g. 100 ms
  - Most threads give up CPU voluntarily much faster

# STCF and STCF-P

- Idea
  - Get the shortest jobs out of the way first
  - Improves short job times significantly
  - Little impact on big ones

## **I. Shortest-Time-to-Completion-First**

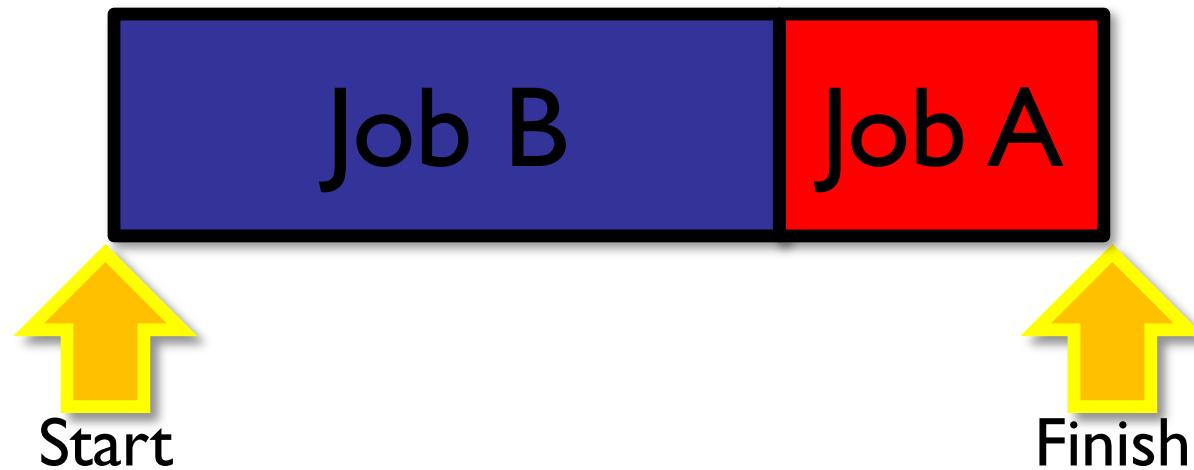
- Run whatever has the least work left before it finishes

## **2. Shortest-Time-to-Completion-First (Pre-emp)**

- If new job arrives with less work than current job has left
- Pre-empt and run the new job

# STCF is optimal

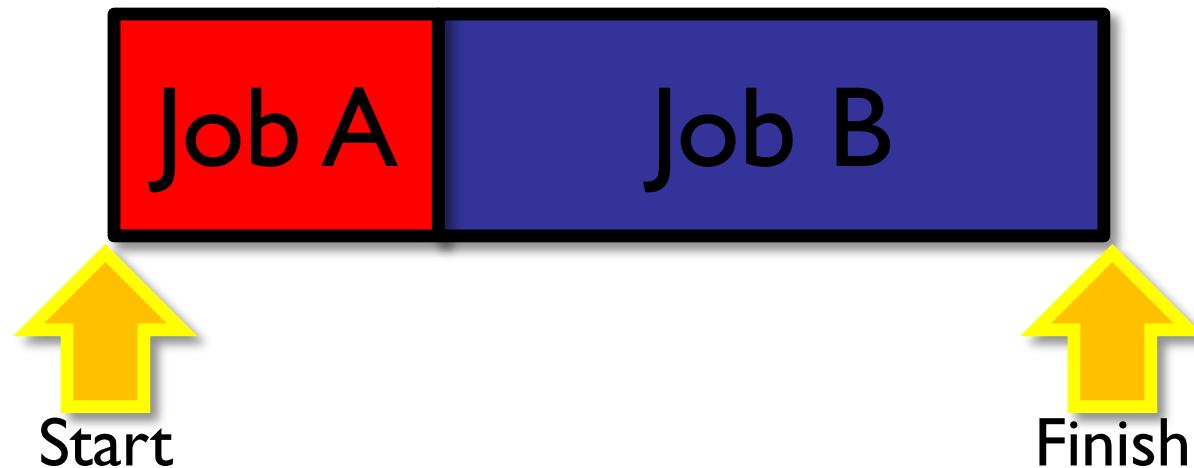
- (among non-pre-emptive policies)



What happened to total time to complete A and B?

# STCF is optimal

- (among non-pre-emptive policies)

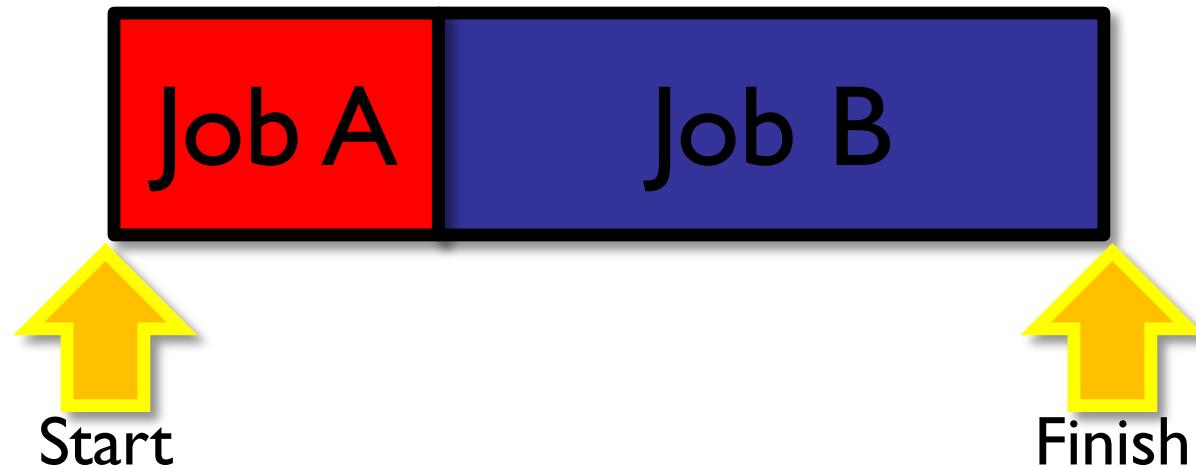


What happened to the time to complete A?

What happened to the time to complete B?

# STCF is optimal

- (among non-pre-emptive policies)

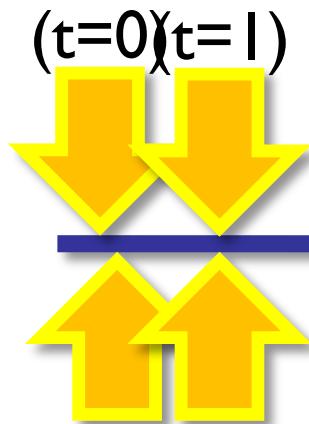


What happened to the average completion time?

# STCF-P is also optimal

- (among pre-emptive policies)
- Job A (100 seconds), Job B (1 second)

A arrives at ~~starts~~  
 $t=0$



A done  
 $t=101$



B arrives ~~and pre-~~  
 $t=1$

~~empt~~  
 $t=0+$

- Average response time = 51 seconds

# STCF-P

- Pros
  - Optimal average response time
- Cons?
  - Can be unfair
  - What happens to long jobs if short jobs keep coming?
  - Legend from the olden days ...
    - IBM 7094 was turned off in 1973
    - Found a “long-running” job from 1967 that hadn’t been scheduled
  - Requires knowledge of the future



# Knowledge of the future

- You will see this a lot.
- Examples?
  - Cache replacement (next reference)
- How do you know how much time jobs take?
  - Ask the user (what if they lie?)
  - Use the past as a predictor of the future

# Grocery store scheduling

- How do grocery stores schedule?
- Kind of like FCFS
- Express lanes
  - Make it kind of like STCF
  - Allow short jobs to get through quickly
- STCF-P would probably be considered unfair



# I/O

- What if a program does I/O too?

```
while (1) {  
    do 1ms of CPU  
    do 10ms of I/O  
}
```

- To scheduler, is this a long job or a short job?
  - Short
  - Thread scheduler only cares about CPU time

# Final example

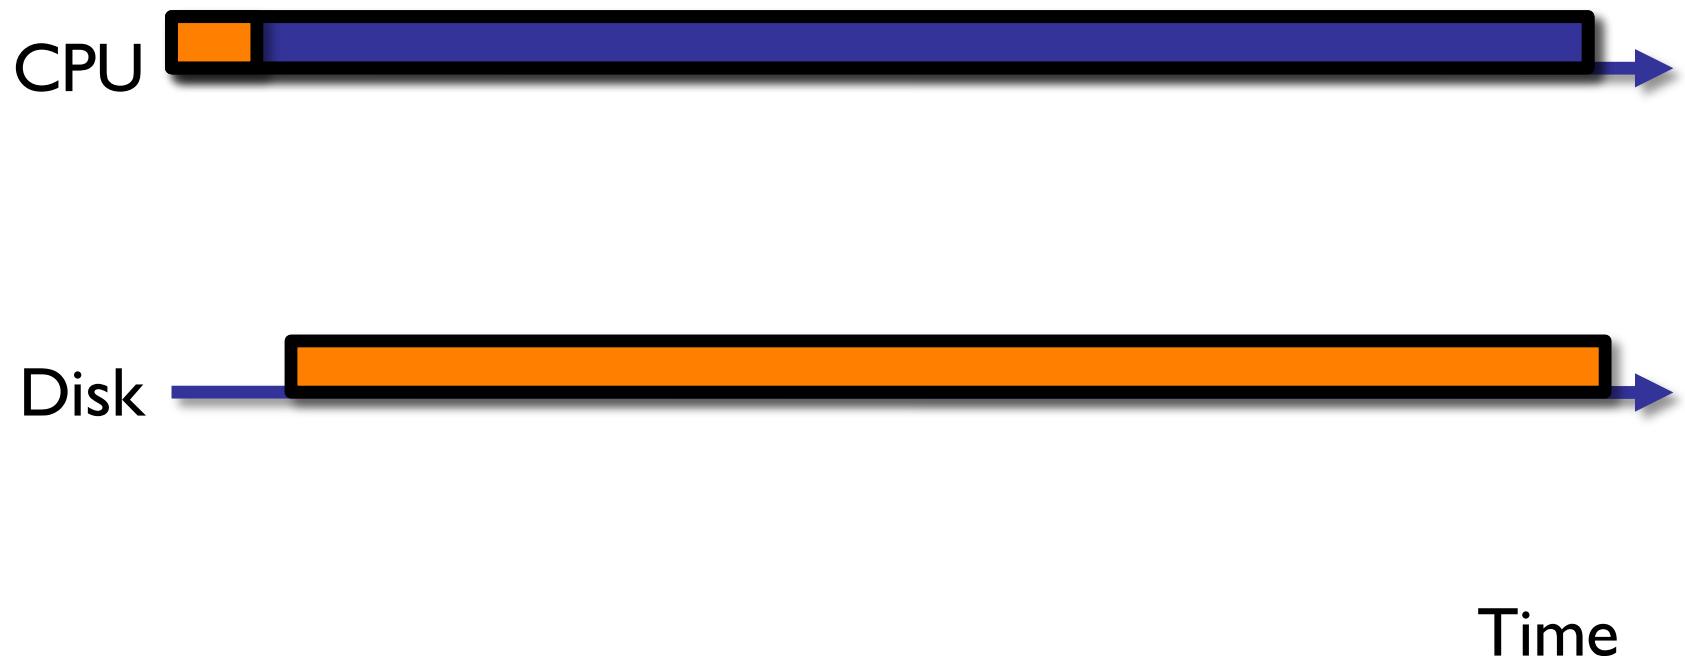
- Job A (CPU-bound)
  - No blocking for I/O, runs for 1000 seconds
- Job B (CPU-bound)
  - No blocking for I/O, runs for 1000 seconds
- Job C (I/O-bound)

```
while (1) {  
    do 1ms of CPU  
    do 10ms of I/O  
}
```

# Each job on its own

A: 100% CPU, 0% Disk   B: 100% CPU, 0% Disk

C: 1/11 CPU, 10/11 Disk



# Mixing jobs (FCFS)

A: 100% CPU, 0% Disk   B: 100% CPU, 0% Disk

C: 1/11 CPU, 10/11 Disk



How well would FCFS work? Not well.

# Mixing jobs (RR with 100ms)

A: 100% CPU, 0% Disk   B: 100% CPU, 0% Disk

C: 1/11 CPU, 10/11 Disk

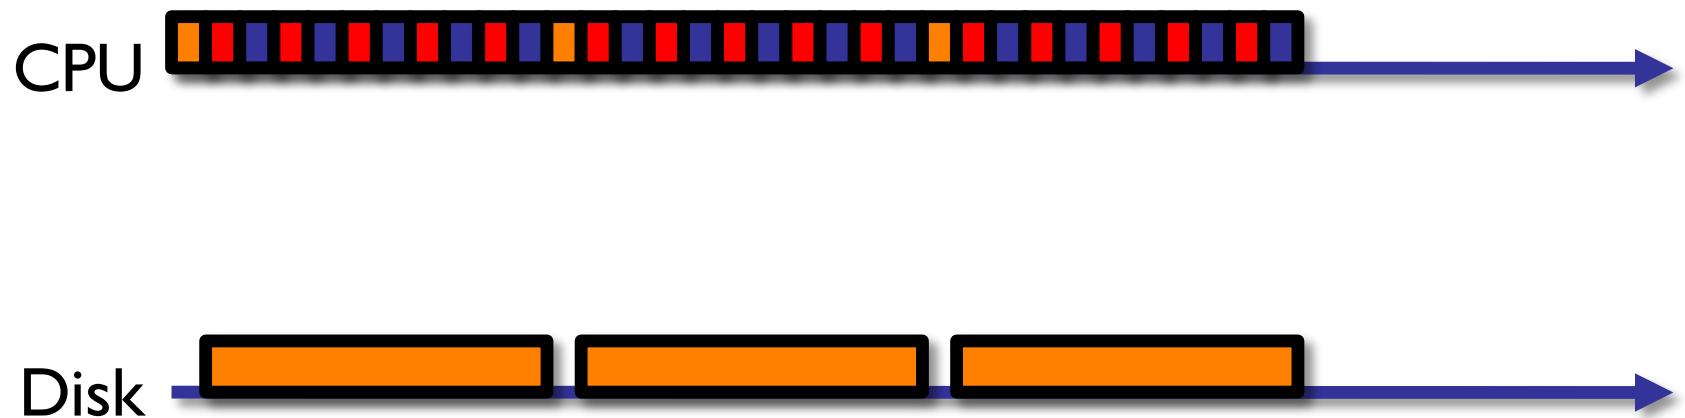


How well would RR with 100 ms slice work?

# Mixing jobs (RR with 1ms)

A: 100% CPU, 0% Disk   B: 100% CPU, 0% Disk

C: 1/11 CPU, 10/11 Disk



How well would RR with 1 ms slice work?

Good Disk utilization (~90%)

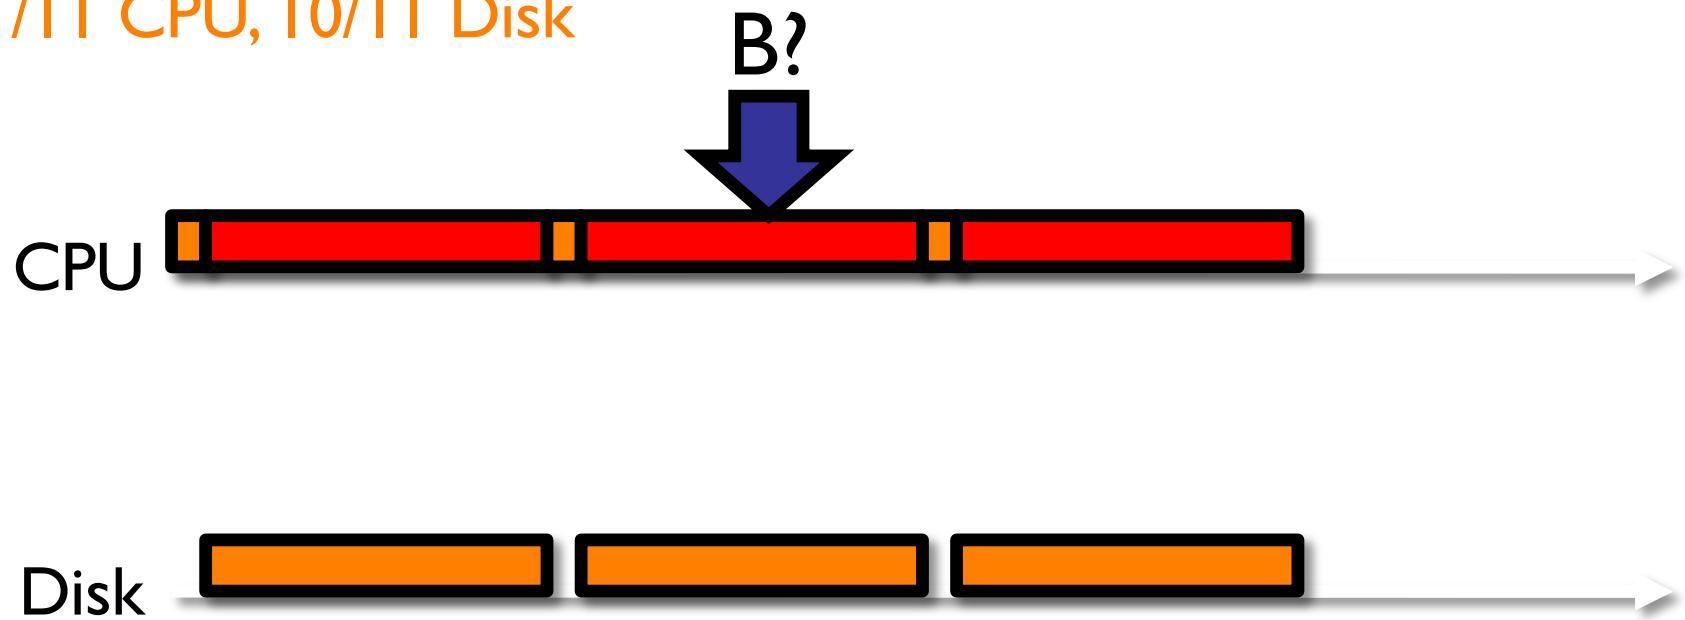
Good principle: start things that can be parallelized early

A lot of context switches though ...

# Mixing jobs (STCF-P)

A: 100% CPU, 0% Disk   B: 100% CPU, 0% Disk

C: 1/I 11 CPU, 10/I 11 Disk



How well would STCF-P work? (run C as soon as its I/O is done)

Good Disk utilization (~90%)

Many fewer context switches

Why not run B here? When will B run? When A finishes

# Real-time scheduling

- So far, we've focused on average-case
- Alternative scheduling goal
  - Finish everything before its deadline
  - Calls for worst-case analysis
- How do we meet all of our deadlines?
  - **Earliest-deadline-first** (optimal)
  - Used by students to complete all homework assignments
  - Used by professors to juggle teaching and research...and life
- Note: sometimes tasks get dropped (for both of us)

# Earliest-deadline-first (EDF)

- EDF
  - Run the job with the earliest deadline
  - If a new job comes in with earlier deadline
    - Pre-empt the current job
    - Start the next one
- This is optimal
  - (assuming it is possible to meet all deadlines)

# EDF example

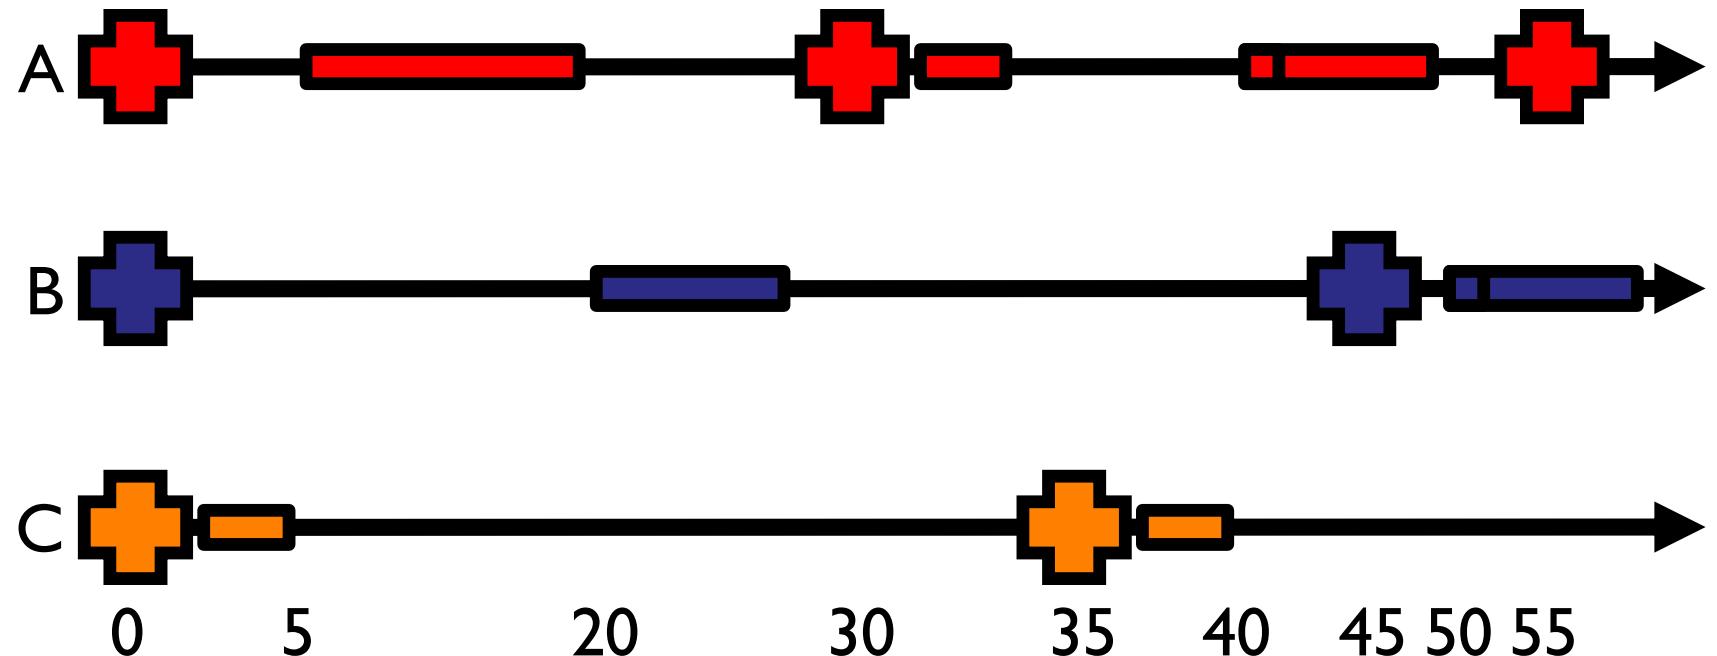
- Job A
  - Takes 15 seconds
  - Due 20 seconds after entry
- Job B
  - Takes 10 seconds
  - Due 30 seconds after entry
- Job C
  - Takes 5 seconds
  - Due 10 seconds after entry

# EDF example

A: takes 15, due in 20

B: takes 10, due in 30

C: takes 5, due in 10



# Proportional-share Schedulers

- A general class of scheduling algorithms
- Process  $P_i$  given a CPU weight  $w_i > 0$
- The scheduler ensures the following
  - *forall*  $i, j$ ,  $|T_i(t_1, t_2)/T_j(t_1, t_2) - w_i/w_j| \leq \epsilon$
  - Given  $P_i$  and  $P_j$  were not blocked during  $[t_1, t_2]$
- Who chooses the weights and how?
  - Application modeling problem: non-trivial
  - Approaches: analytical, empirical

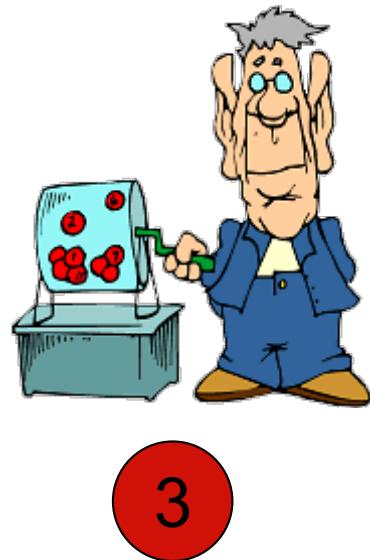
# Lottery Scheduling

- Perhaps the simplest proportional-share scheduler
- Create lottery tickets equal to the sum of the weights of all processes
- Draw a lottery ticket and schedule the process that owns that ticket
  - What if the process is not ready?
  - Draw tickets only for ready processes

# Lottery Scheduling

$P1=6$

$P2=9$



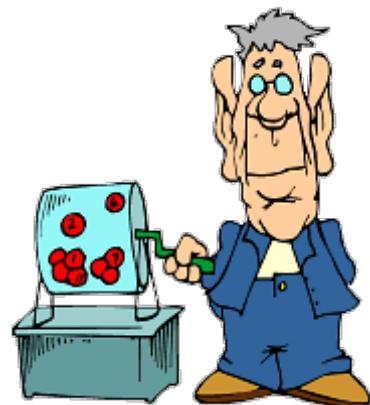
1	4	7	10	13
2	5	8	11	14
3	6	9	12	15

*Schedule P1*

# Lottery Scheduling

$P1=6$

$P2=9$



9

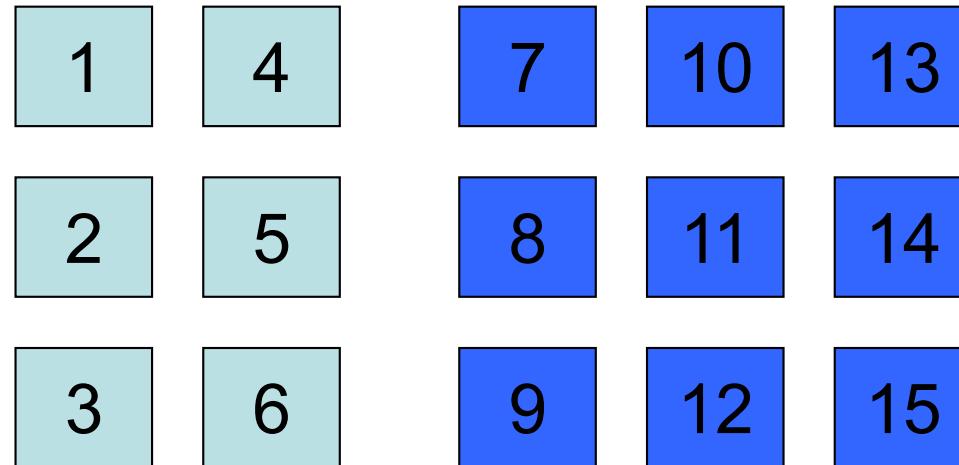
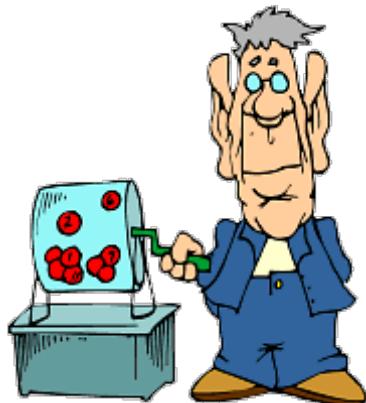
1	4	7	10	13
2	5	8	11	14
3	6	9	12	15

*Schedule P2*

# Lottery Scheduling

$P1=6$

$P2=9$



11

*Schedule P2*

- As  $t$  approaches infinite, processes will get their share (unless they were blocked a lot)
  - $P1 = 6/15 = 40\%$ , and  $P2 = 9/15 = 60\%$
- Problem: only a probabilistic guarantee
- What does the scheduler have to do
  - When a new process arrives?
  - When a process terminates?

# Lottery Scheduling

- **Simple Implementation of proportional-share**
  - Used by many VMs to fairly share CPU of a single computer
  - Implementation is efficient and requires no history
  - Currency abstraction is hierarchical
    - Can delegate lottery tickets to threads or other processes
  - What happens if I take lottery tickets away after you “spend” them?
- **Proportional-share is important in other areas**
  - Example: bandwidth allocation, memory
  - Conflicts between principals (e.g., VMs) may reduce efficiency
    - As with a mechanical disk